

# Kubernetes Architecture

---





# Module Outline

Big picture view

Masters

Nodes

Pods

Services

Deployments


Recap

# Kubernetes Big Picture View

Team




Manager  
(coach)



Team


<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>




Team




Manager  
(coach)







HTTPS




HTTPS




Search




Auth




K/V store





MySQL




Log




Load  
balancer








Nodes a.k.a. Minions




Nodes a.k.a. Minions

in charge!

do the work!



Deployment



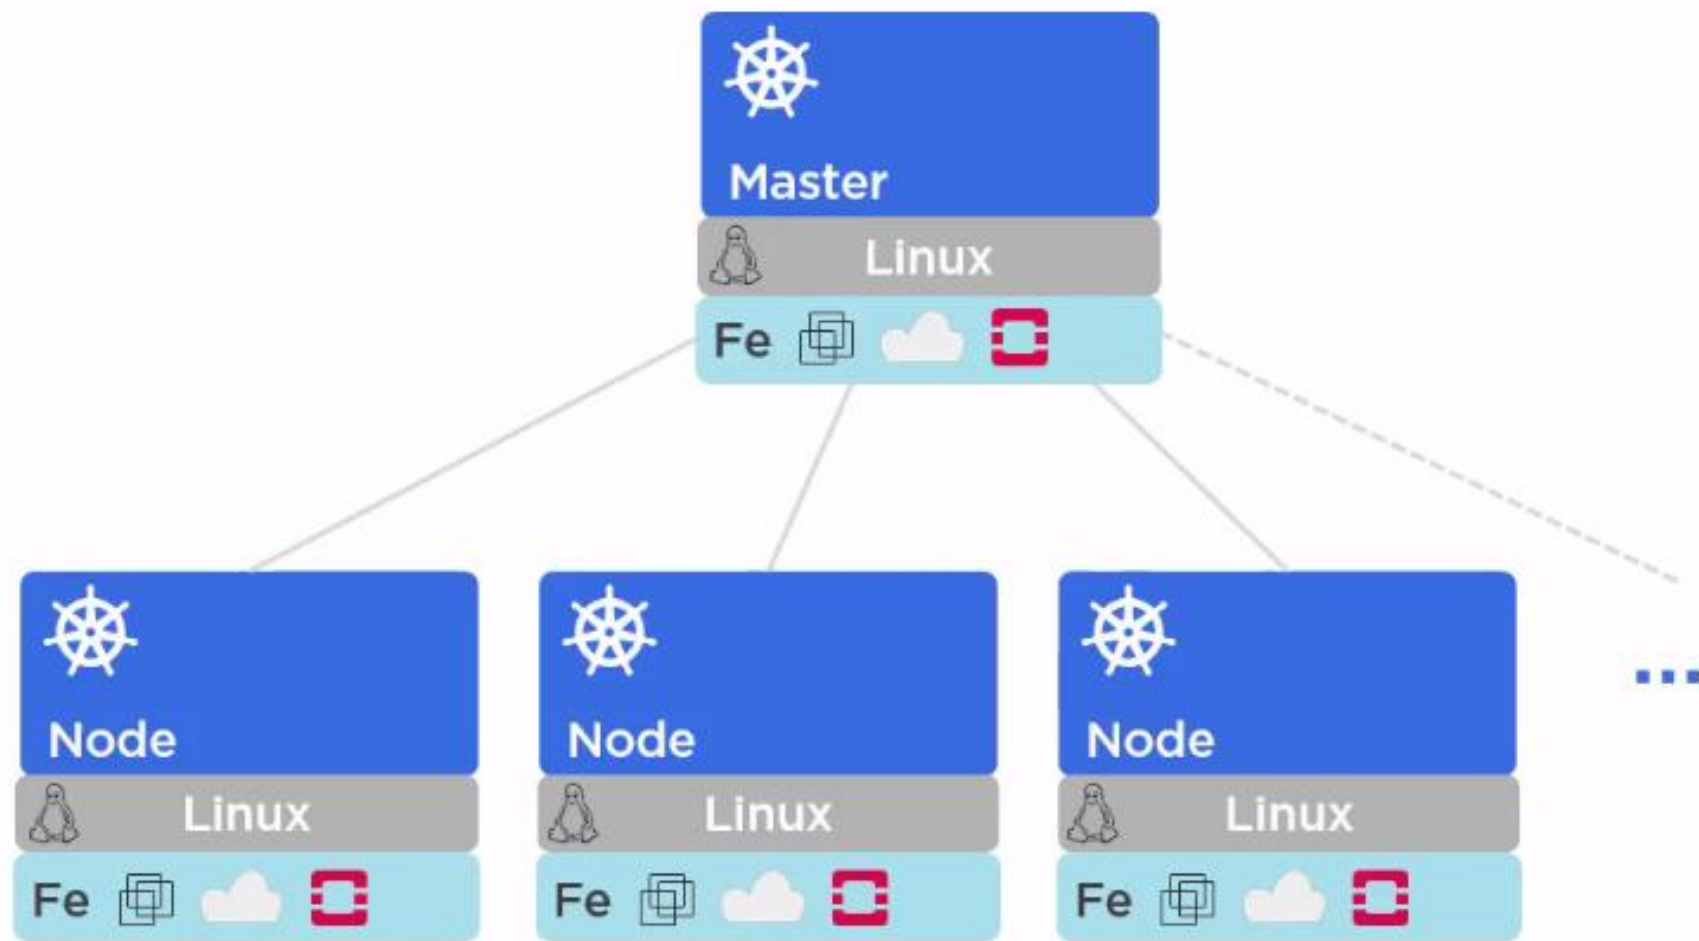
K8s cluster




Nodes a.k.a. Minions

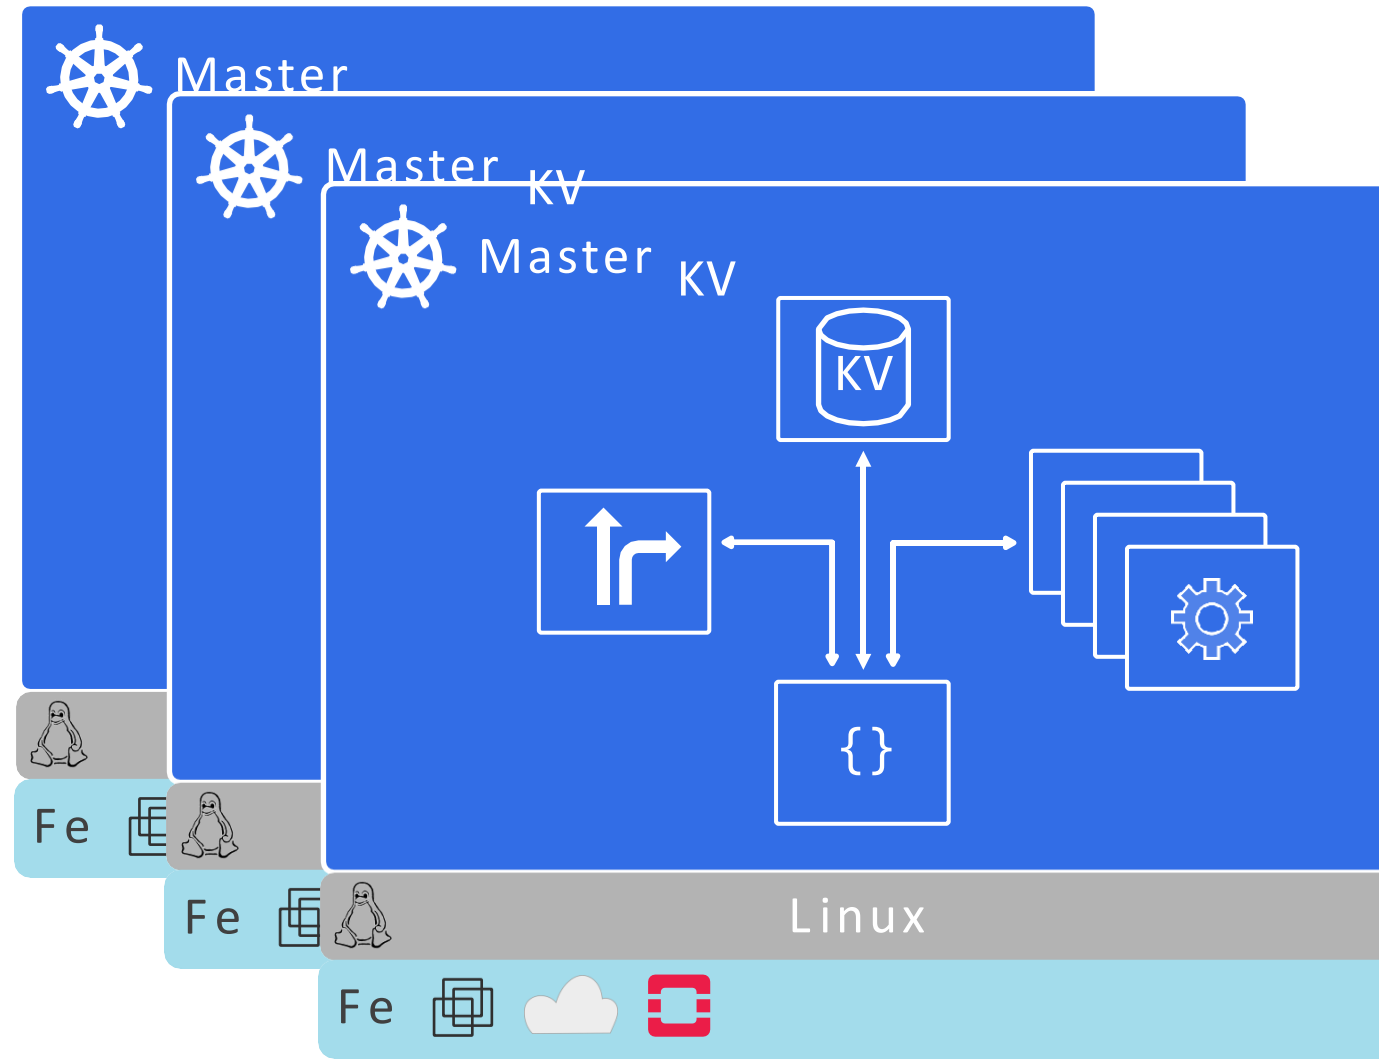
# Masters




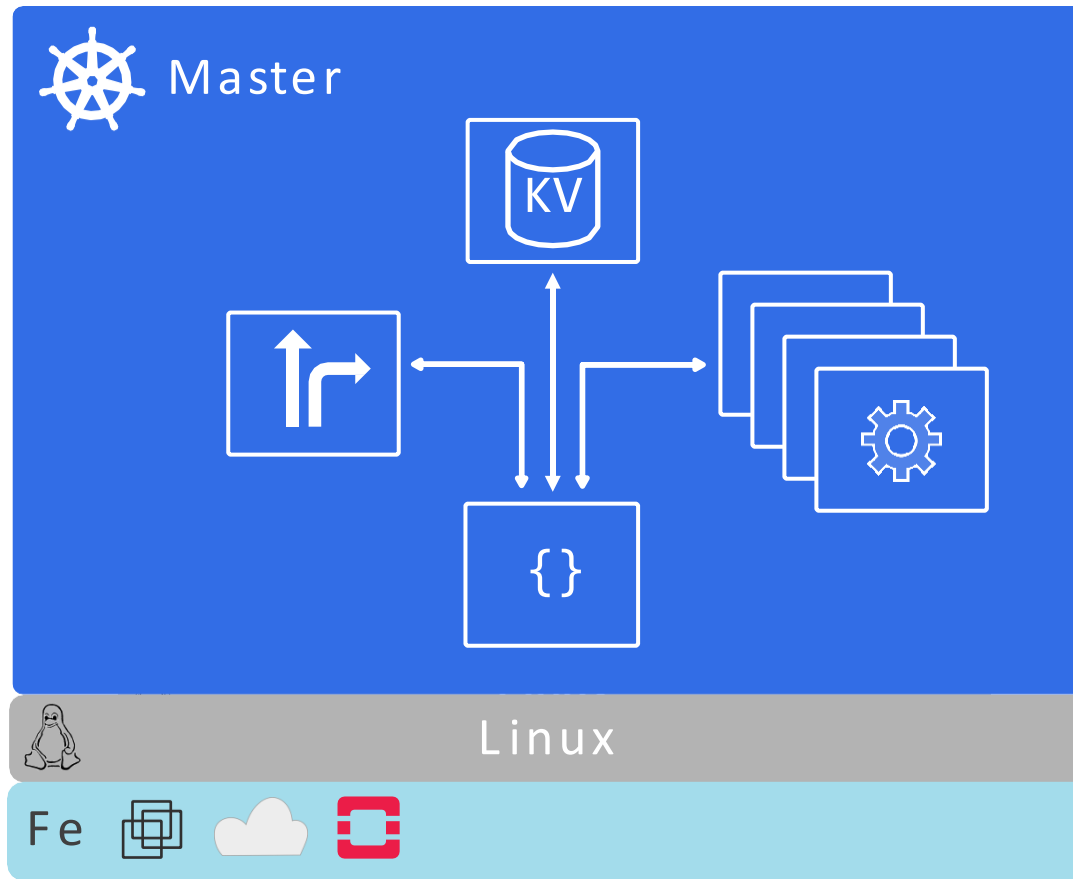

The Kubernetes Control Plane






# Multi-master HA





Don't run user workloads on  
"Master"






# kube-apiserver

Front-end to the control plane

Exposes the API (REST)

Consumes JSON  
(via manifest files)



# Cluster store

Persistent storage


Cluster state and config

Uses etcd

Distributed, consistent,  
watchable...

The "*source of truth*" for  
the cluster

Have a backup plan for it!



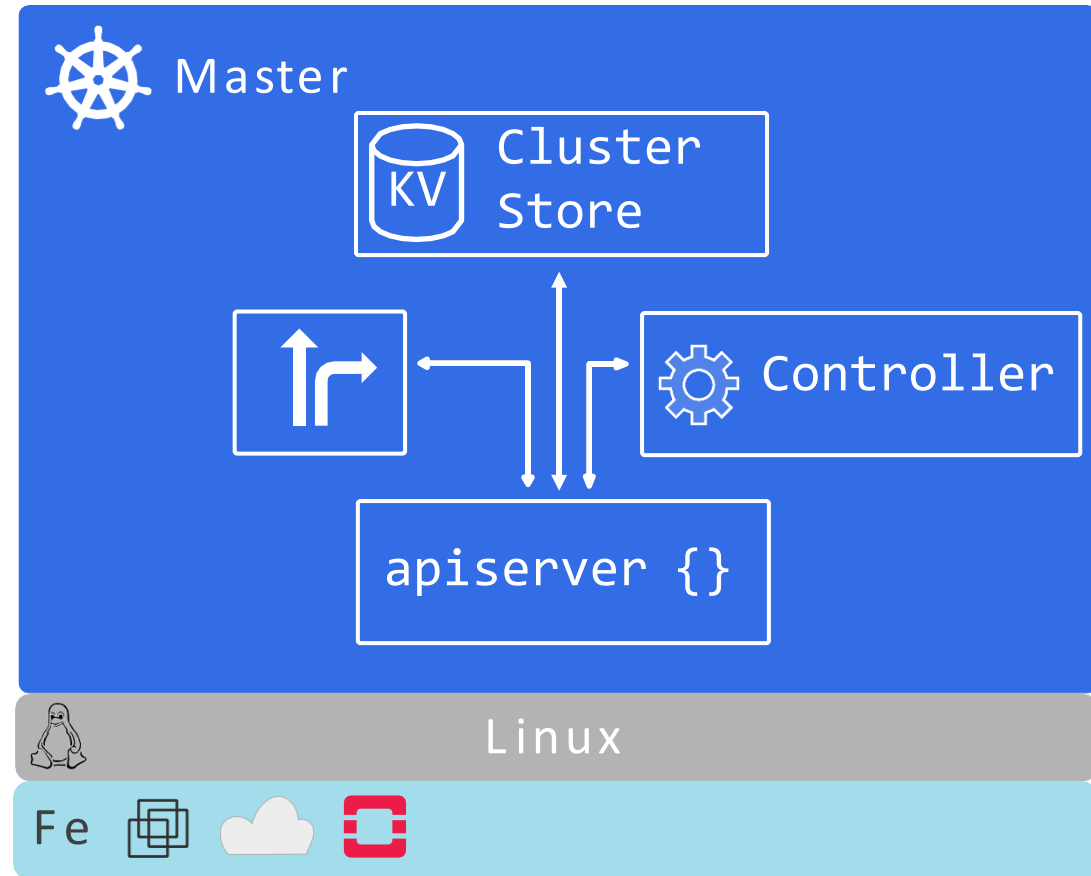
# kube-controller-manager

Controller of controllers

- Node controller
- Endpoints controller
- Namespace controller
- ...

Watches for changes

Helps maintain *desired state*

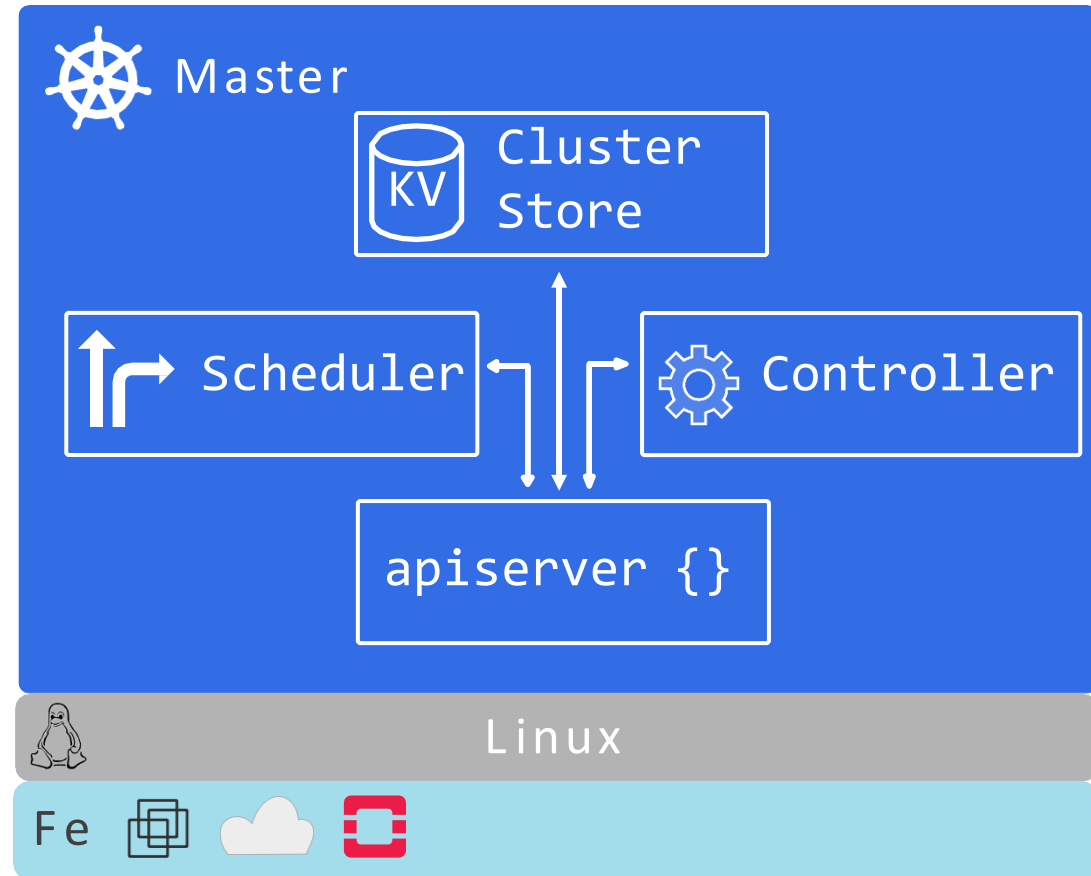



# kube-scheduler

Watches apiserver for new pods


Assigns work to nodes

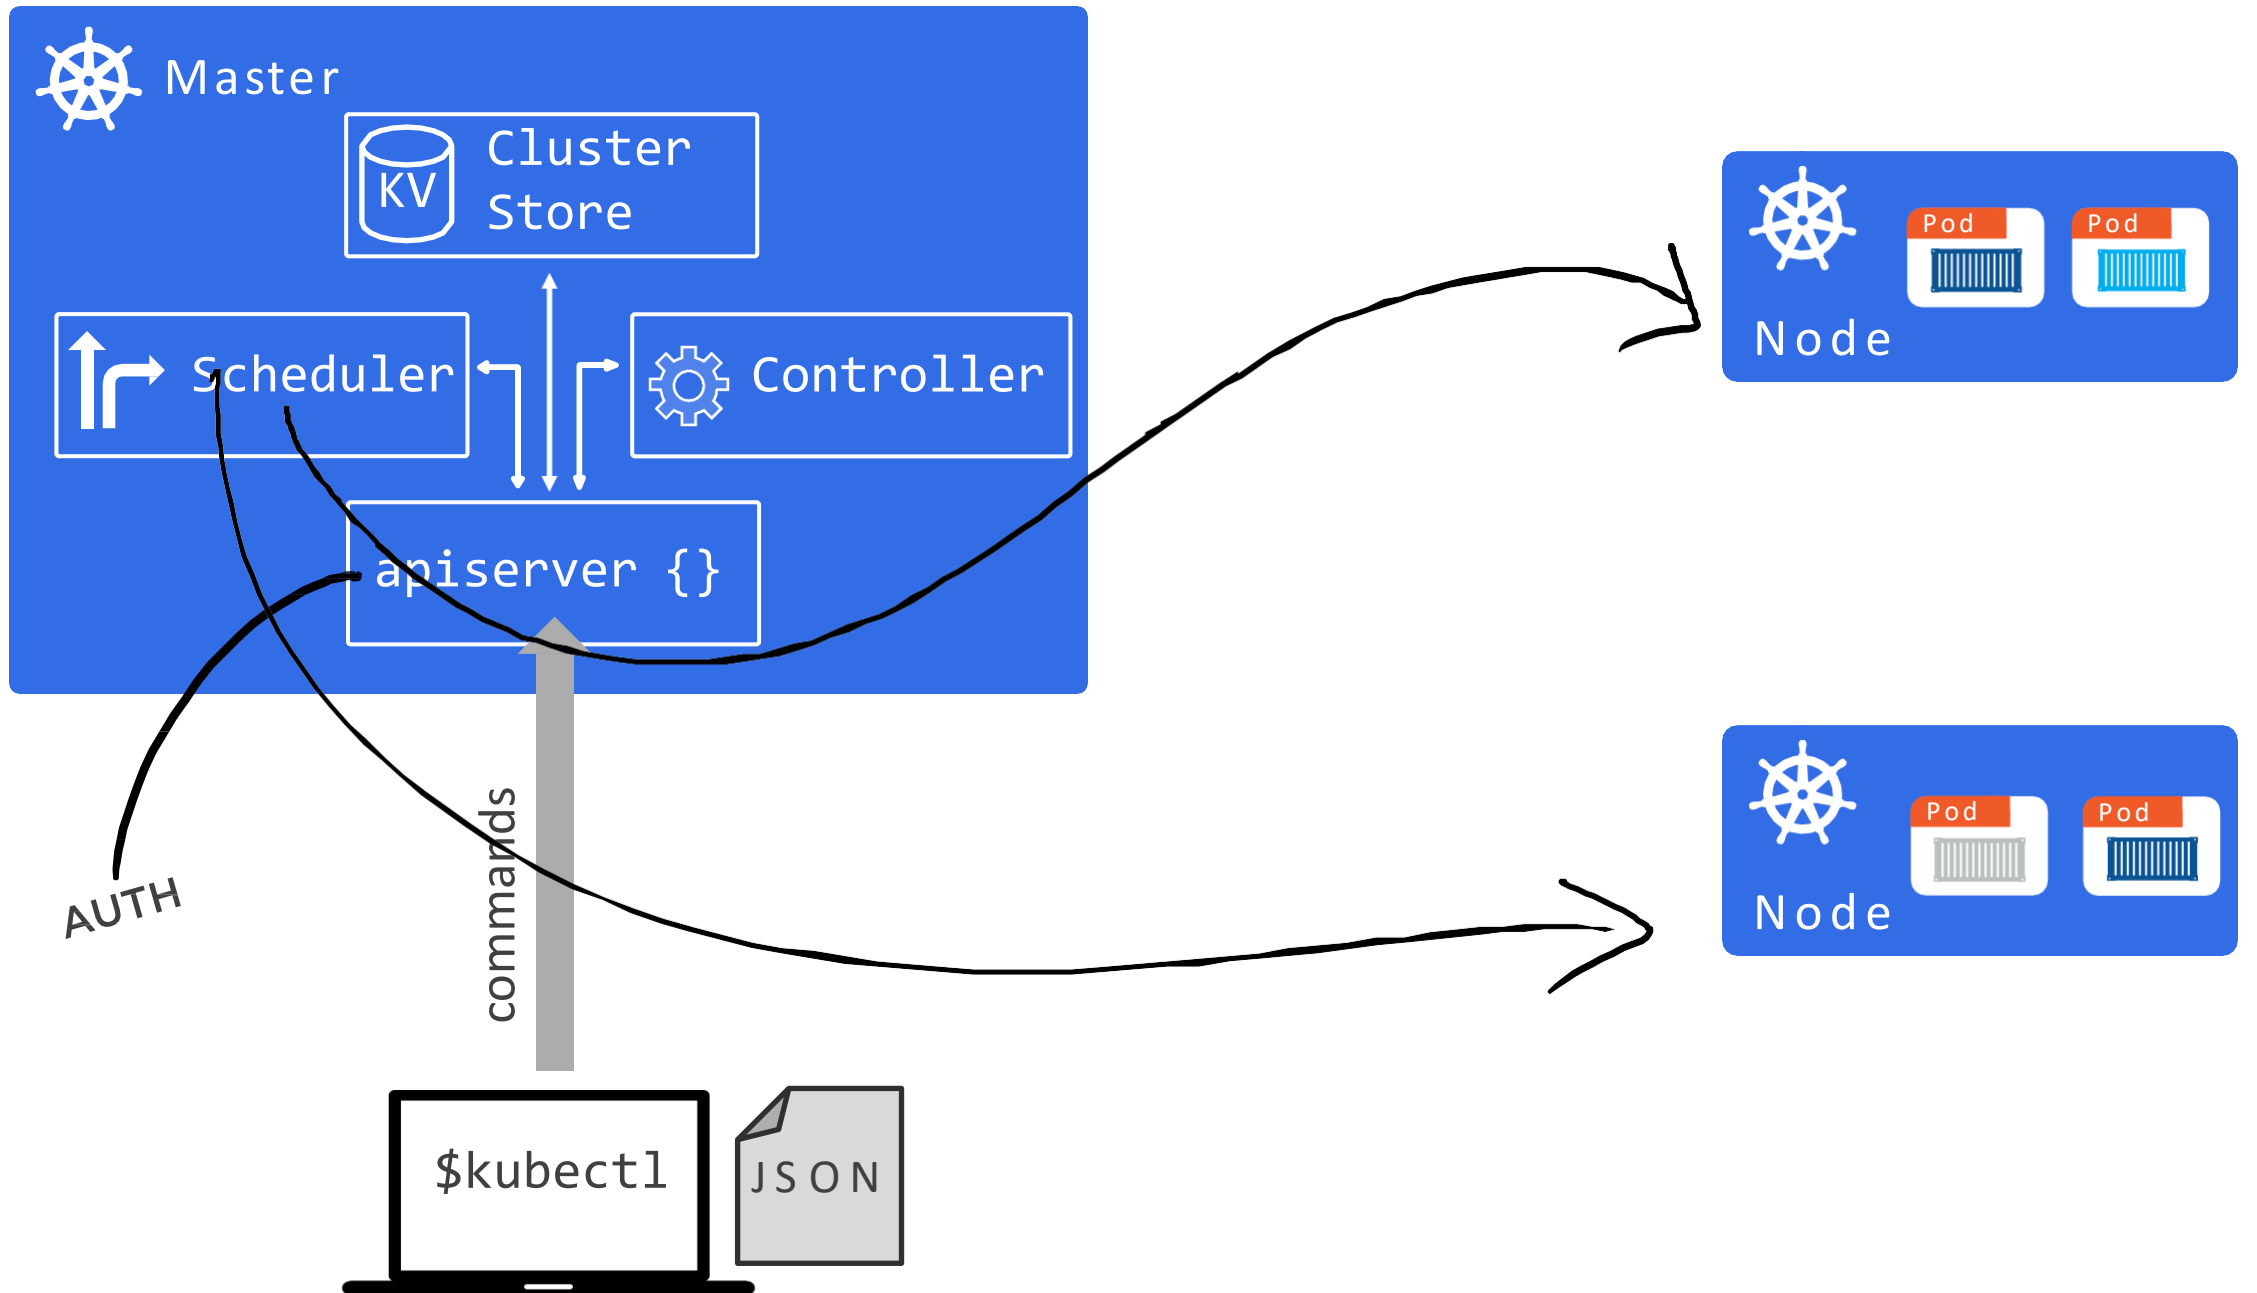
- affinity/anti-affinity
- constraints
- resources
- ...





a.k.a  
master








# Nodes a.k.a “Minions”

The Kubernetes Worker



Node



Linux


Fe






## Kubelet

- The main Kubernetes agent
- Registers node with cluster
- Watches apiserver
- Instantiates pods
- Reports back to master
- Exposes endpoint on :10255






## Container Engine

Does container management:

- Pulling images
- Starting/stopping containers
- ...

Pluggable:

- Usually Docker
- Can be rkt




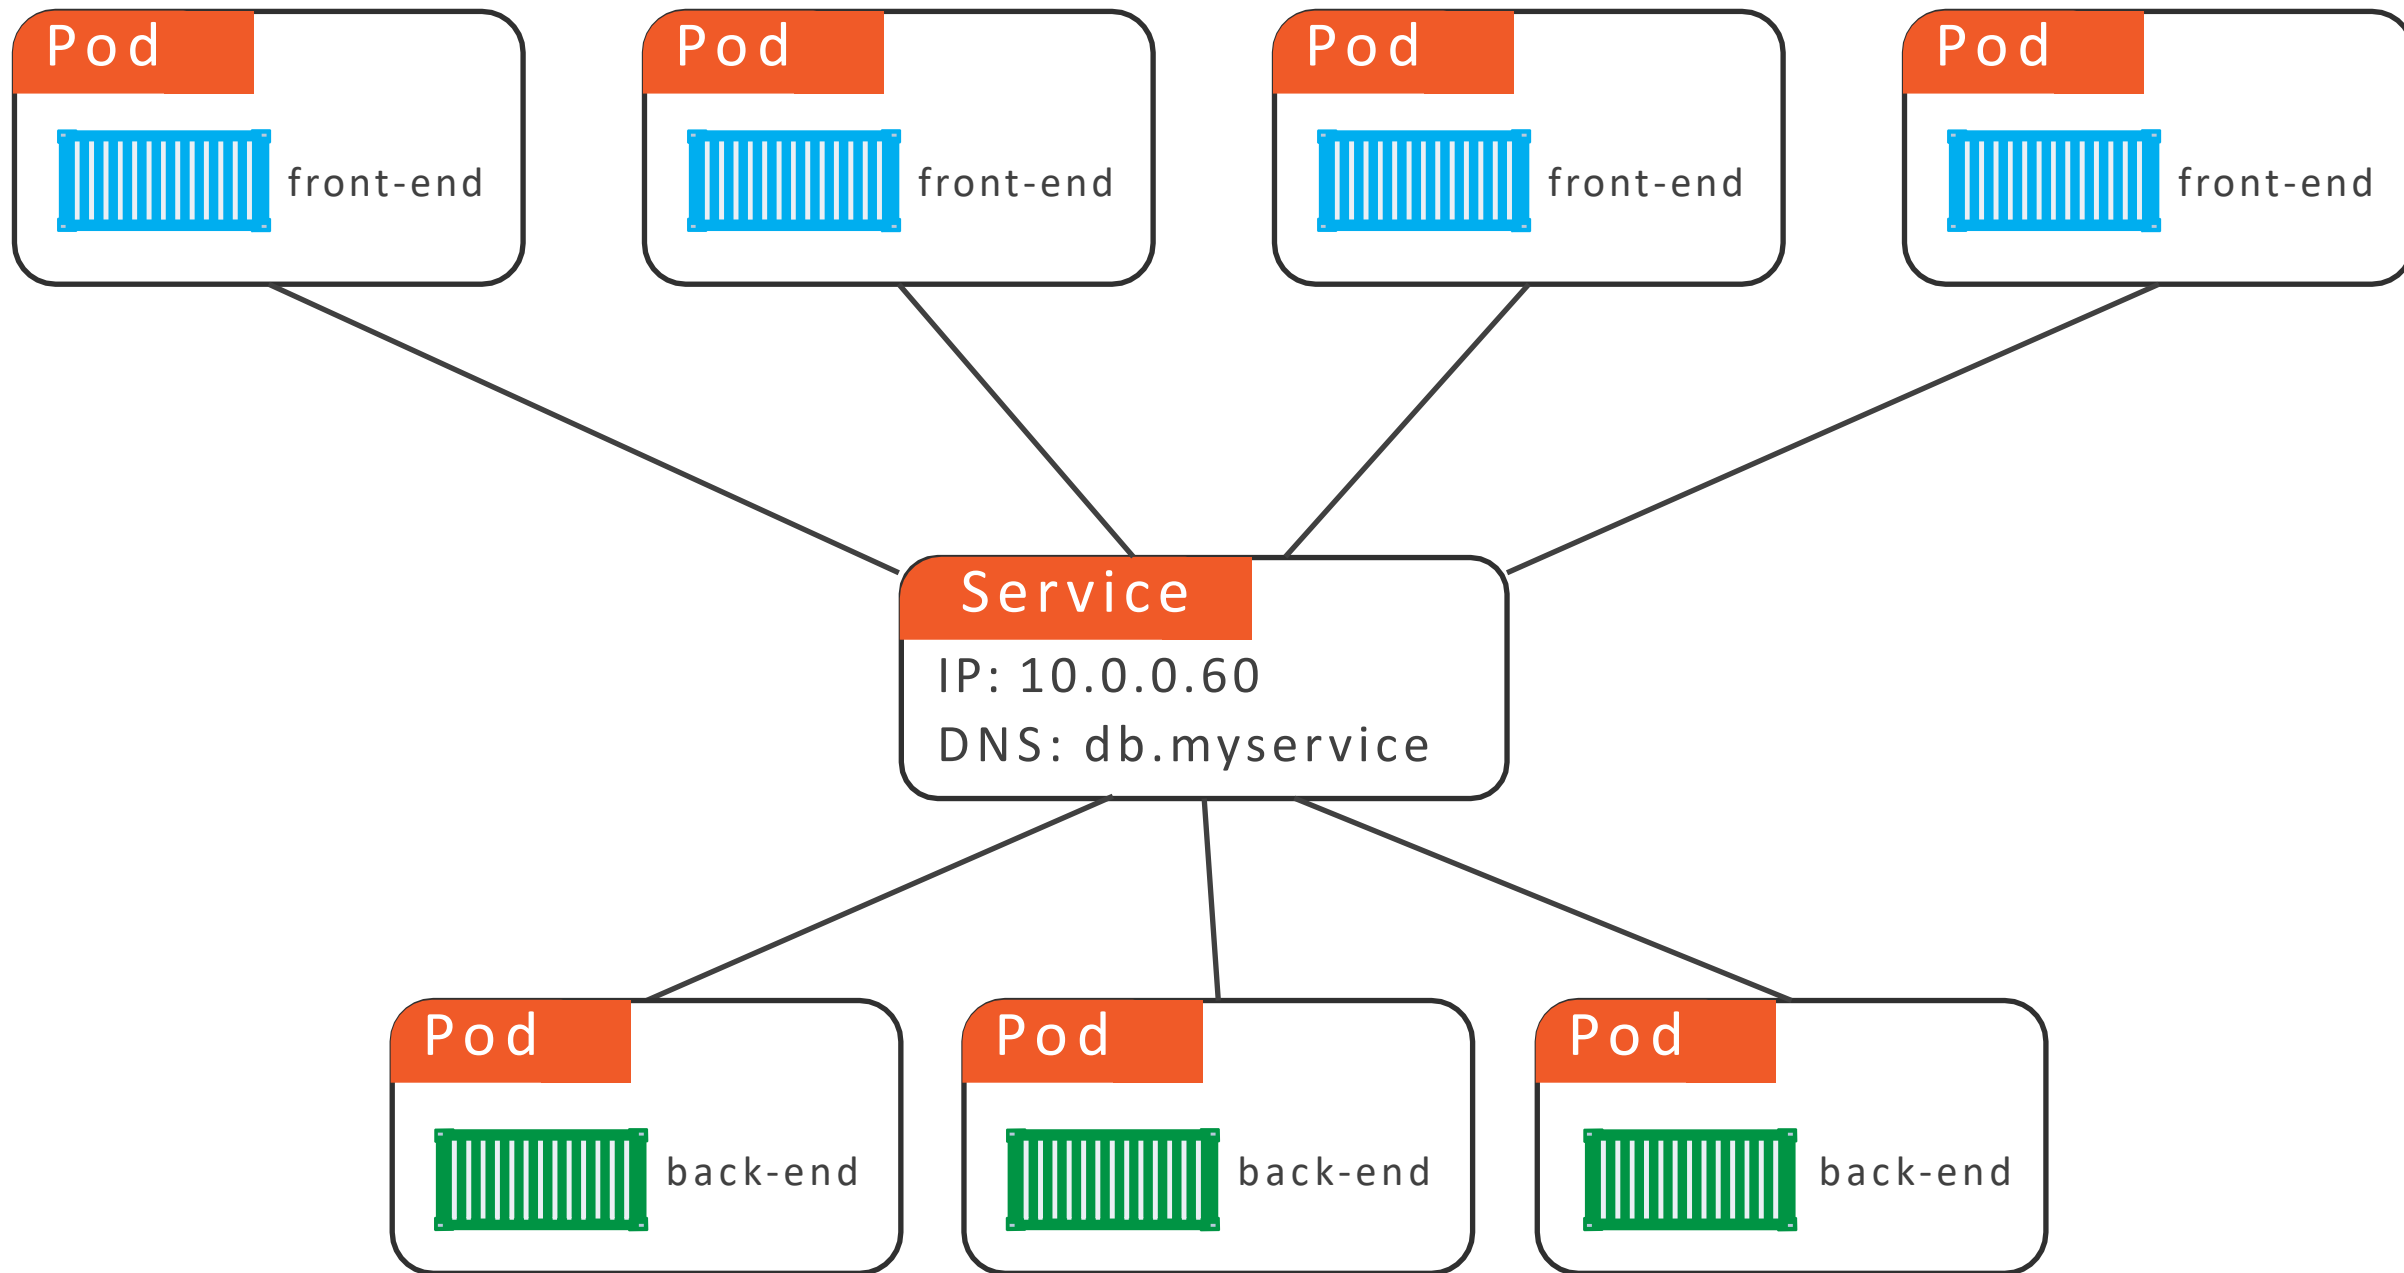


# kube-proxy

Kubernetes networking:


- Pod IP addresses
  - All containers in a pod share a single IP
- Load balances across all pods in a service








Kubelet  
Main Kubernetes agent



Container engine  
Docker or rkt




kube-proxy  
Kubernetes networking




# Declarative Model & Desired State





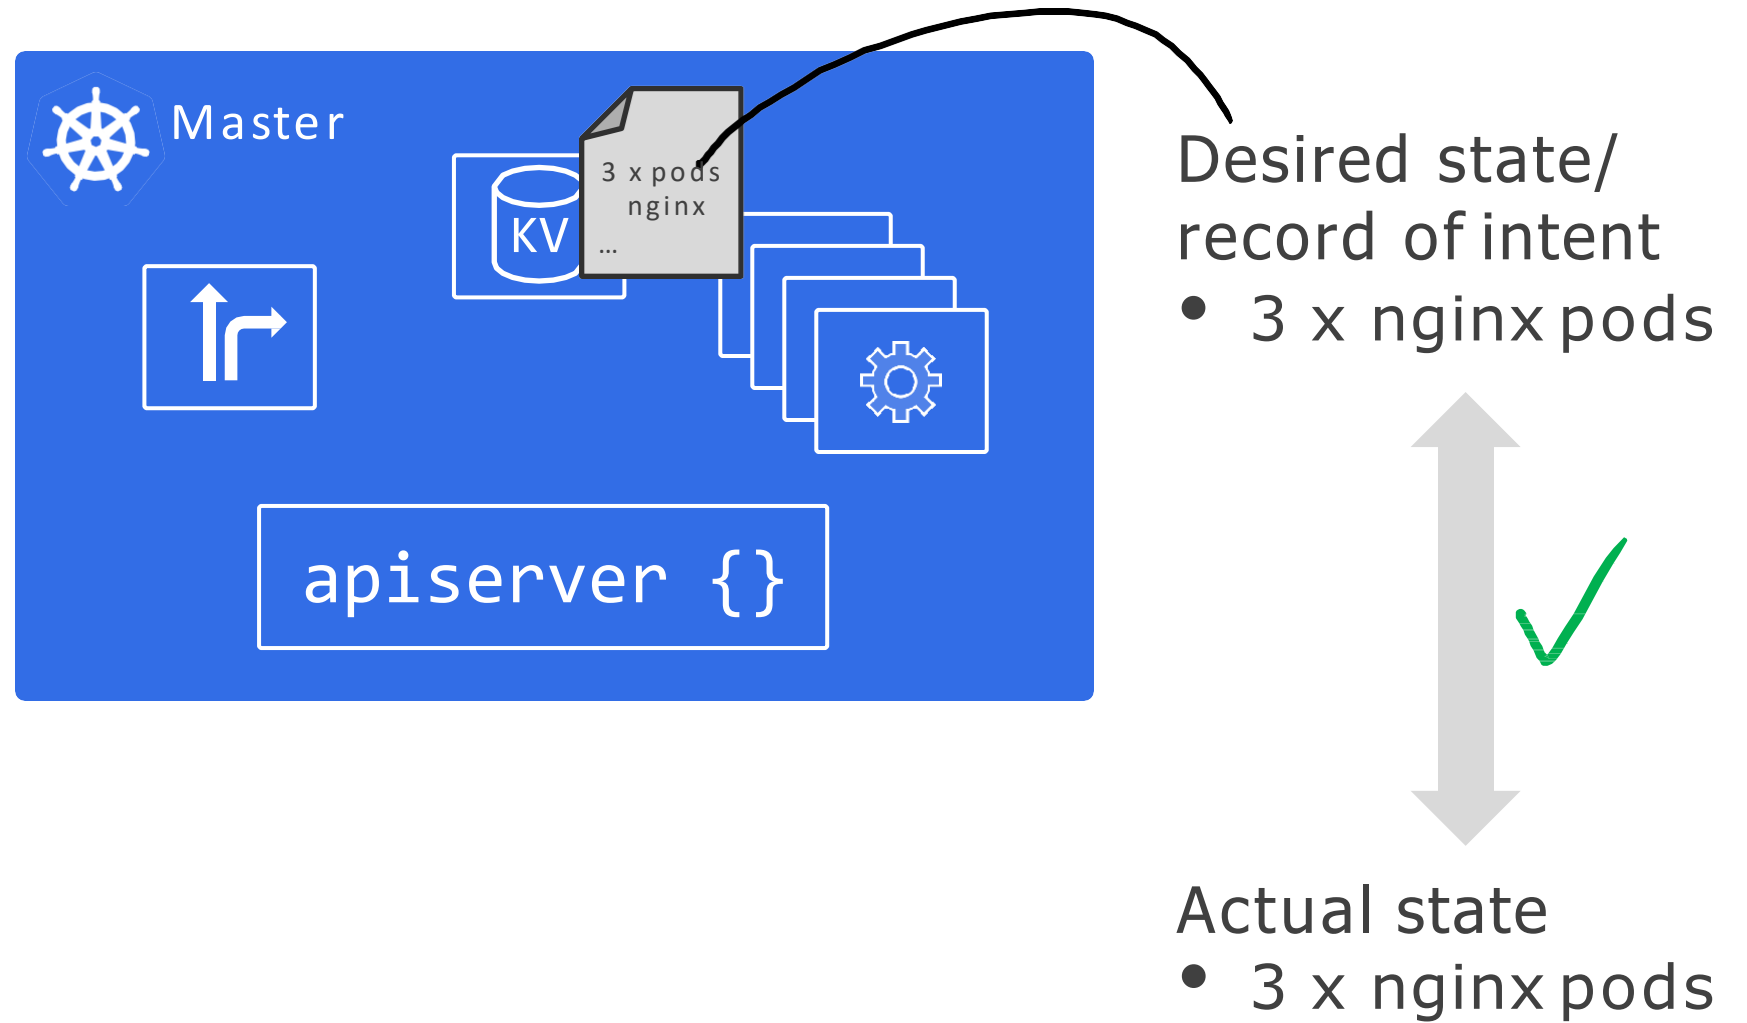



Manifest  
file

YAML or JSON  
Describe desired  
state











Desired state/  
record of intent


- 3 x nginx pods



Actual state


- 2 x nginx pods






Desired state/  
record of intent


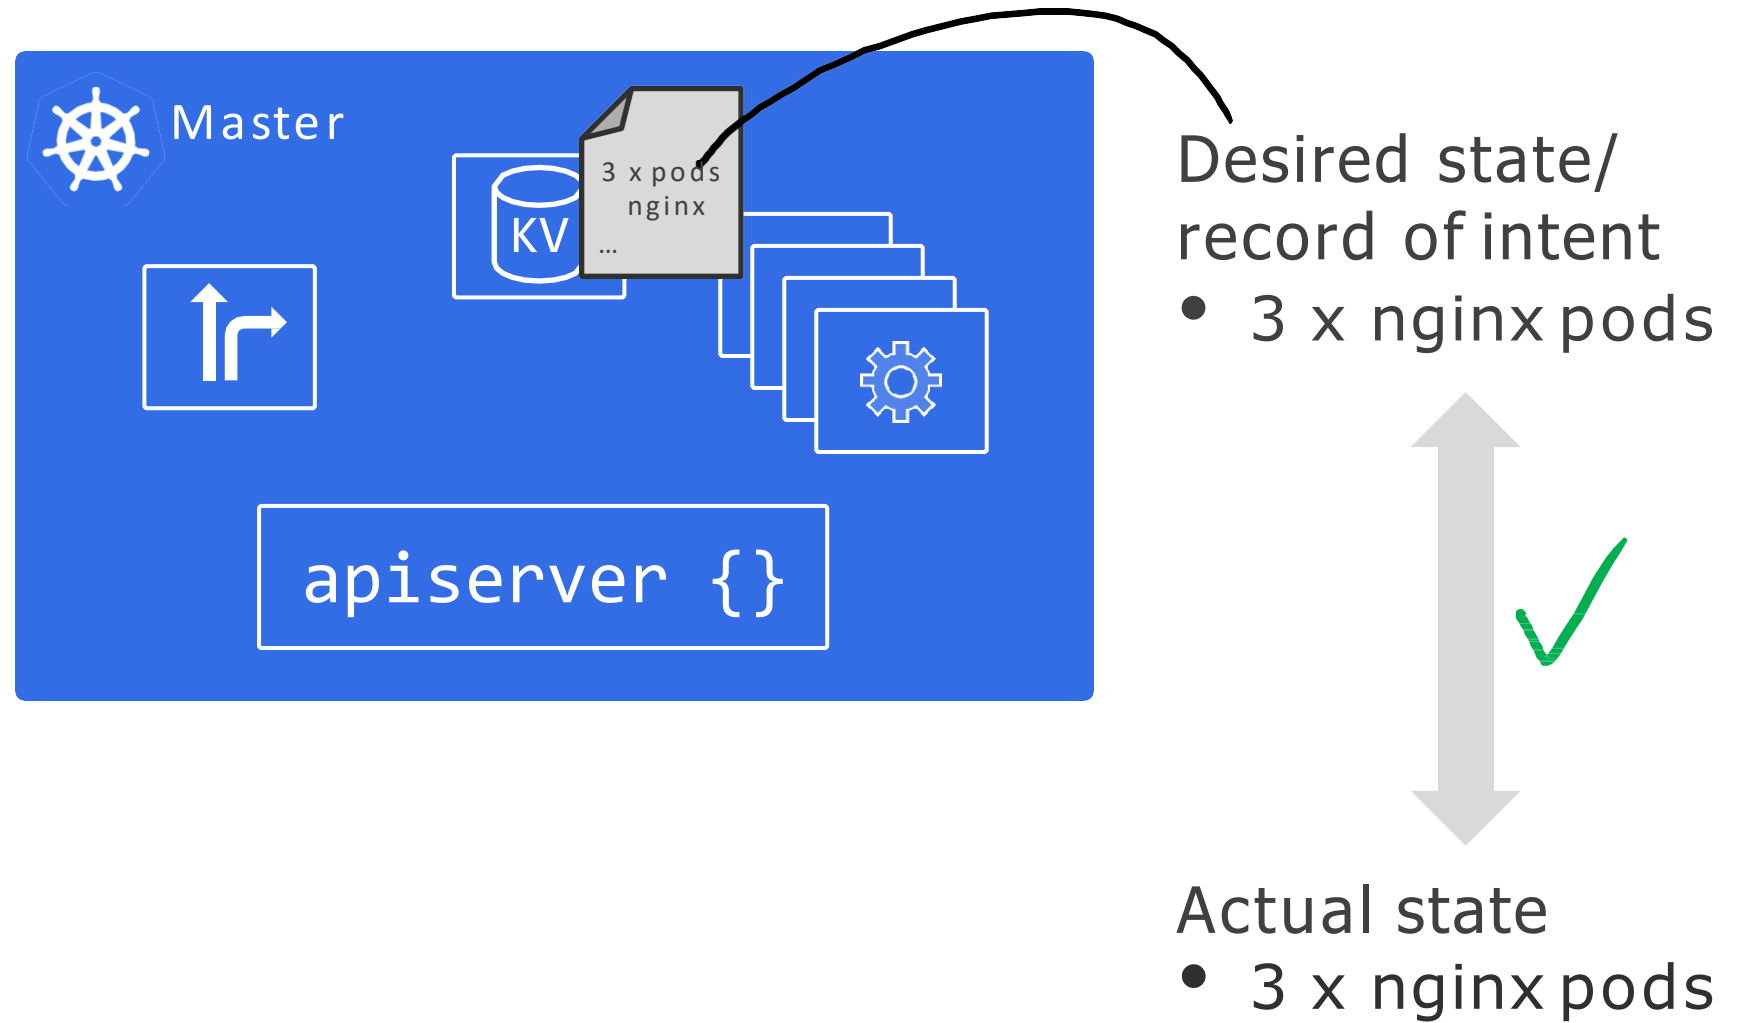
- 3 x nginx pods



Actual state

- 2 x nginx pods








Pods





VM



Container




Pod

Atomic units of scheduling



Containers always run  
inside of pods

Pods can have multiple  
containers  
(advanced use-case)

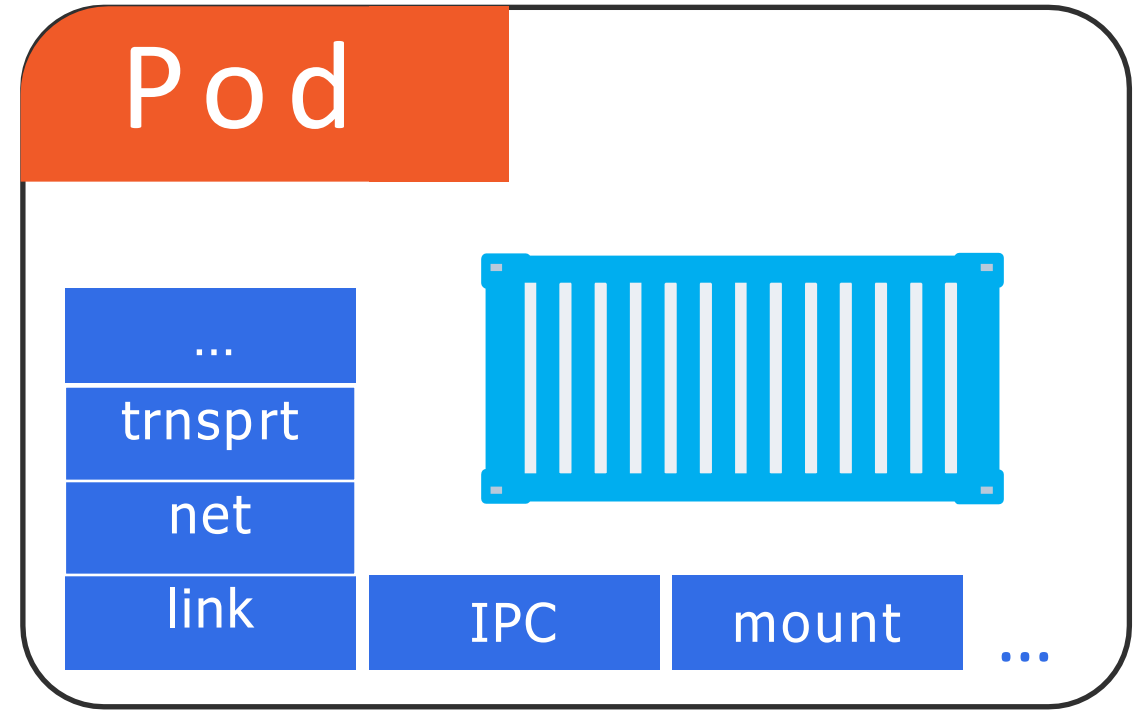


## Ring-fenced environment


- Network stack
- Kernel namespaces
- ...

$n$  containers



All containers in pod share the pod environment






# Tight Coupling




# Loose Coupling



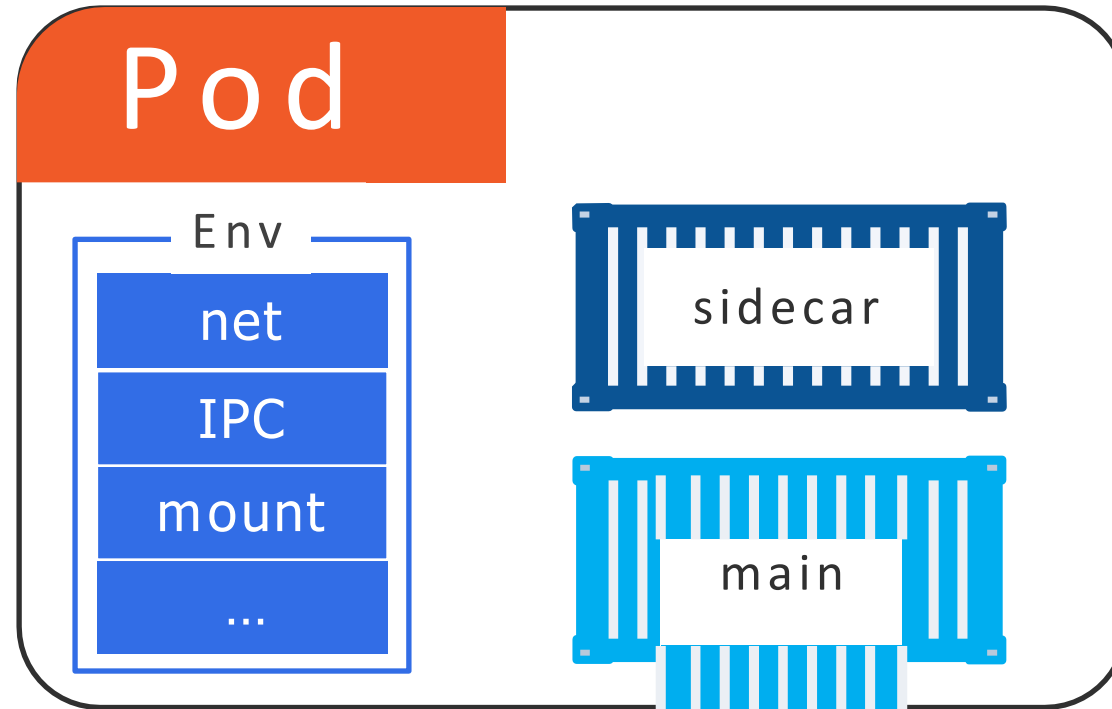
# Pods and Scaling



# Multi-container Pods




# Pods are Atomic

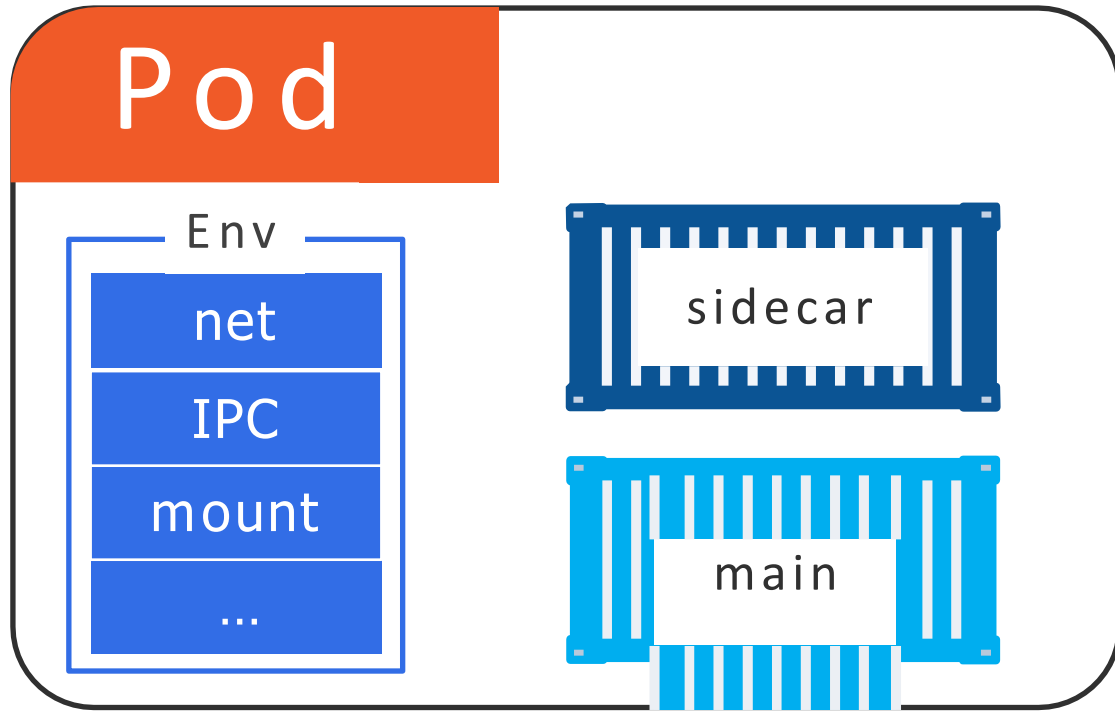




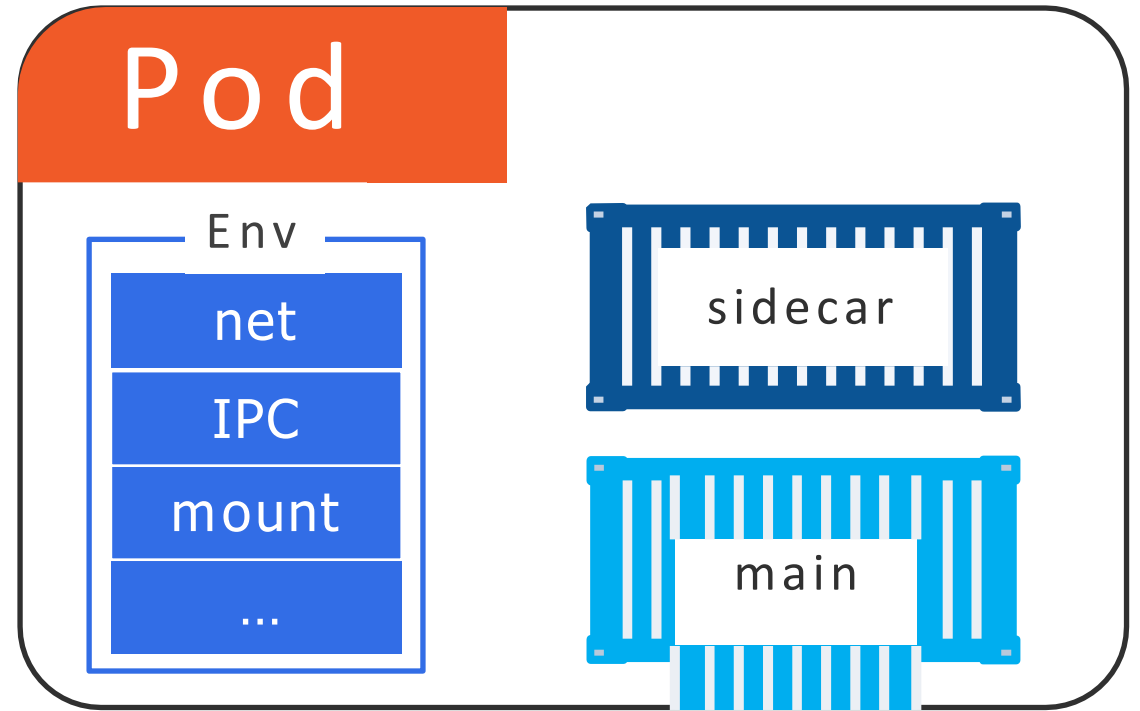
# Pods are Atomic



# Pods are Atomic

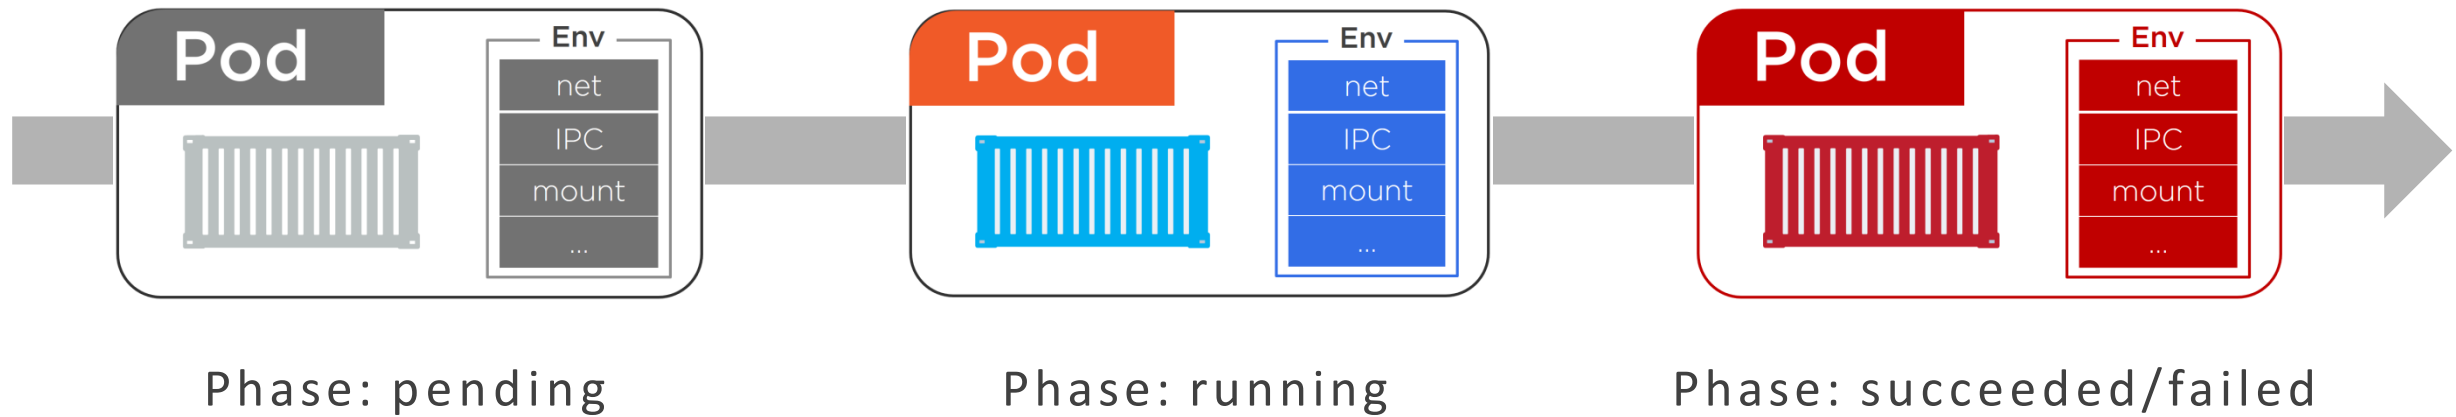


#1 Status:ready





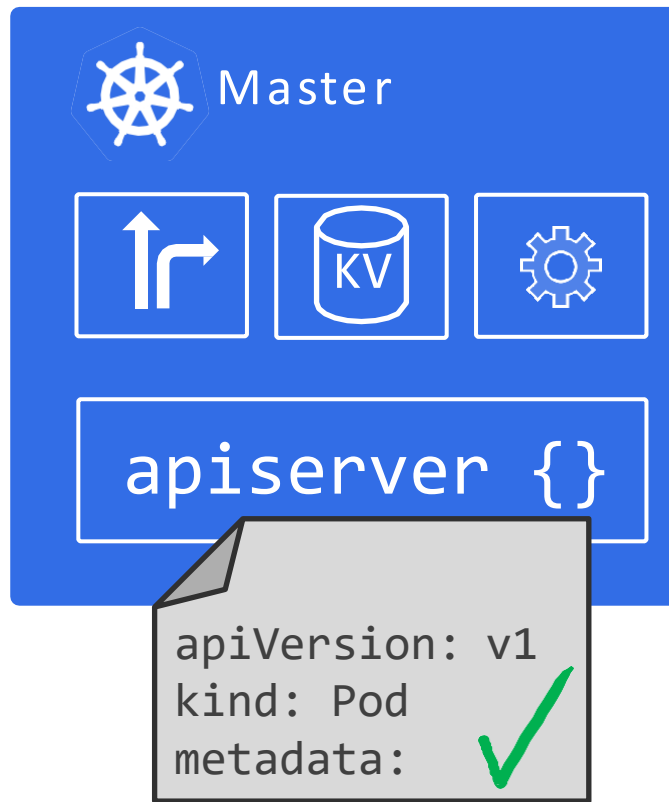
#2 Status:pending

# Pod Lifecycle



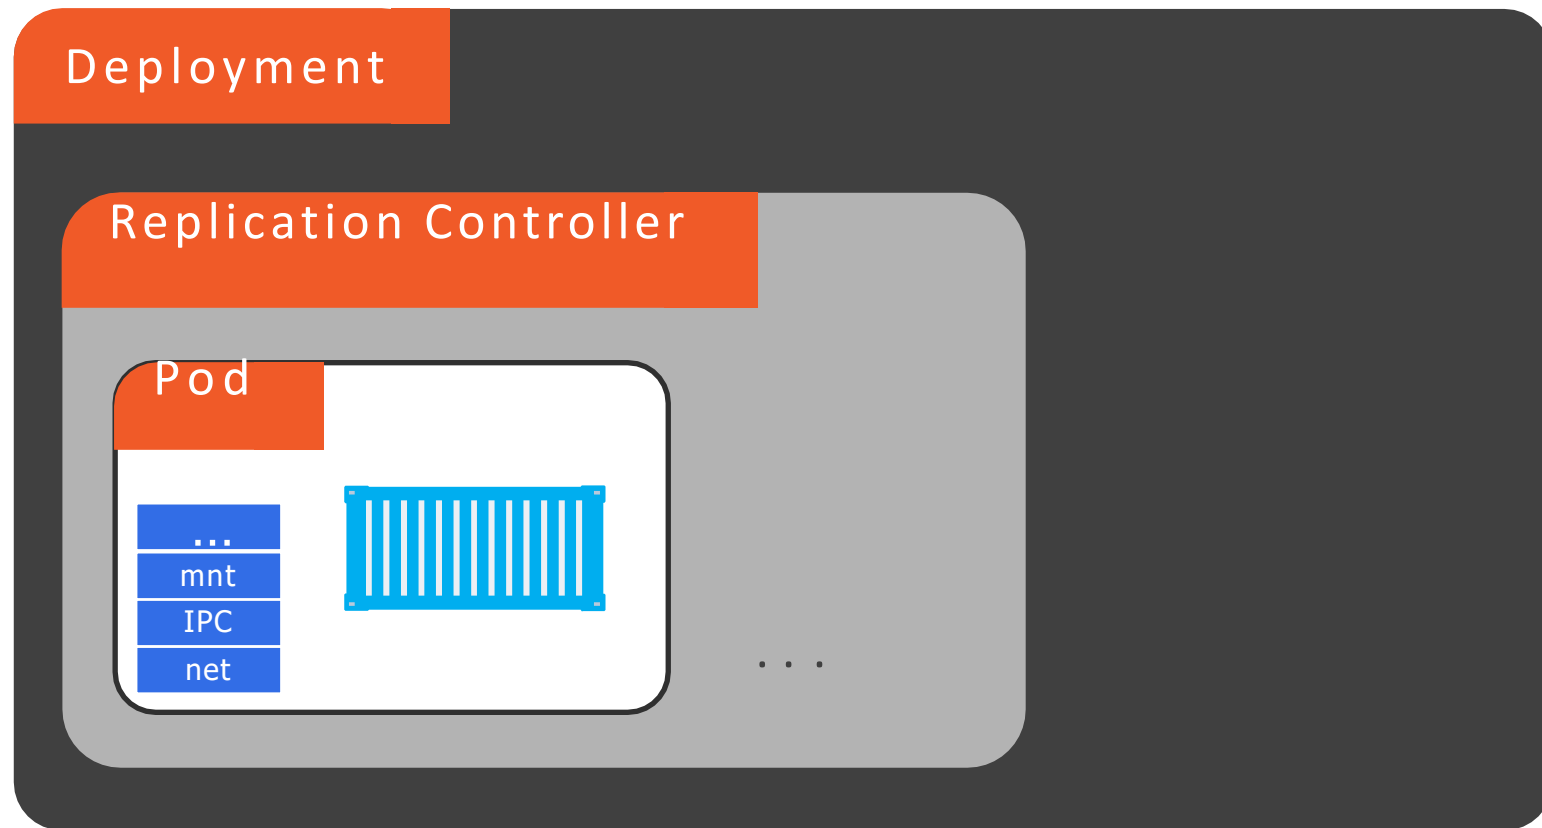
# Deploying Pods

Usually via higher level objects



# Deploying Pods

Usually via higher level objects

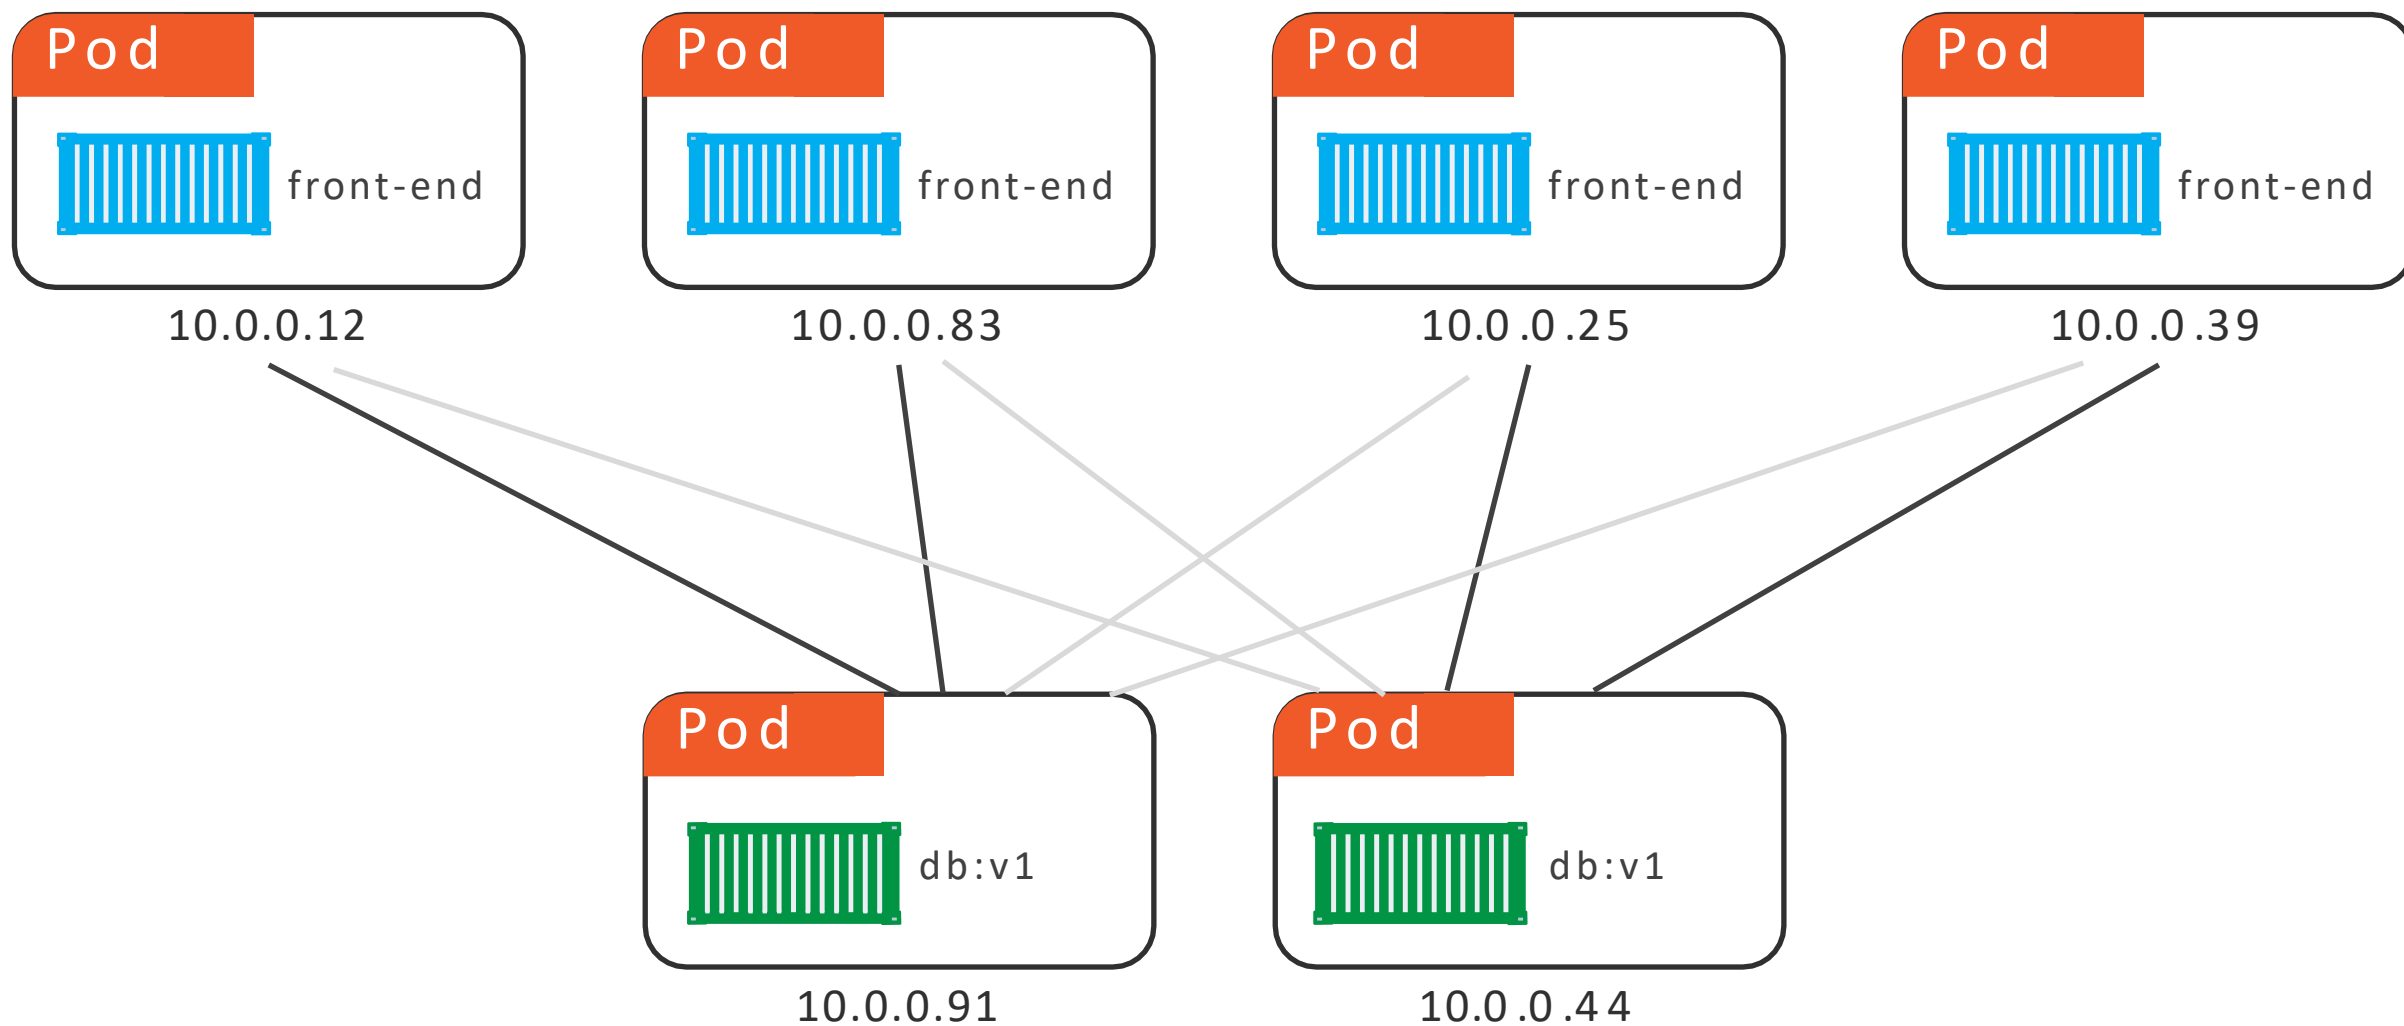


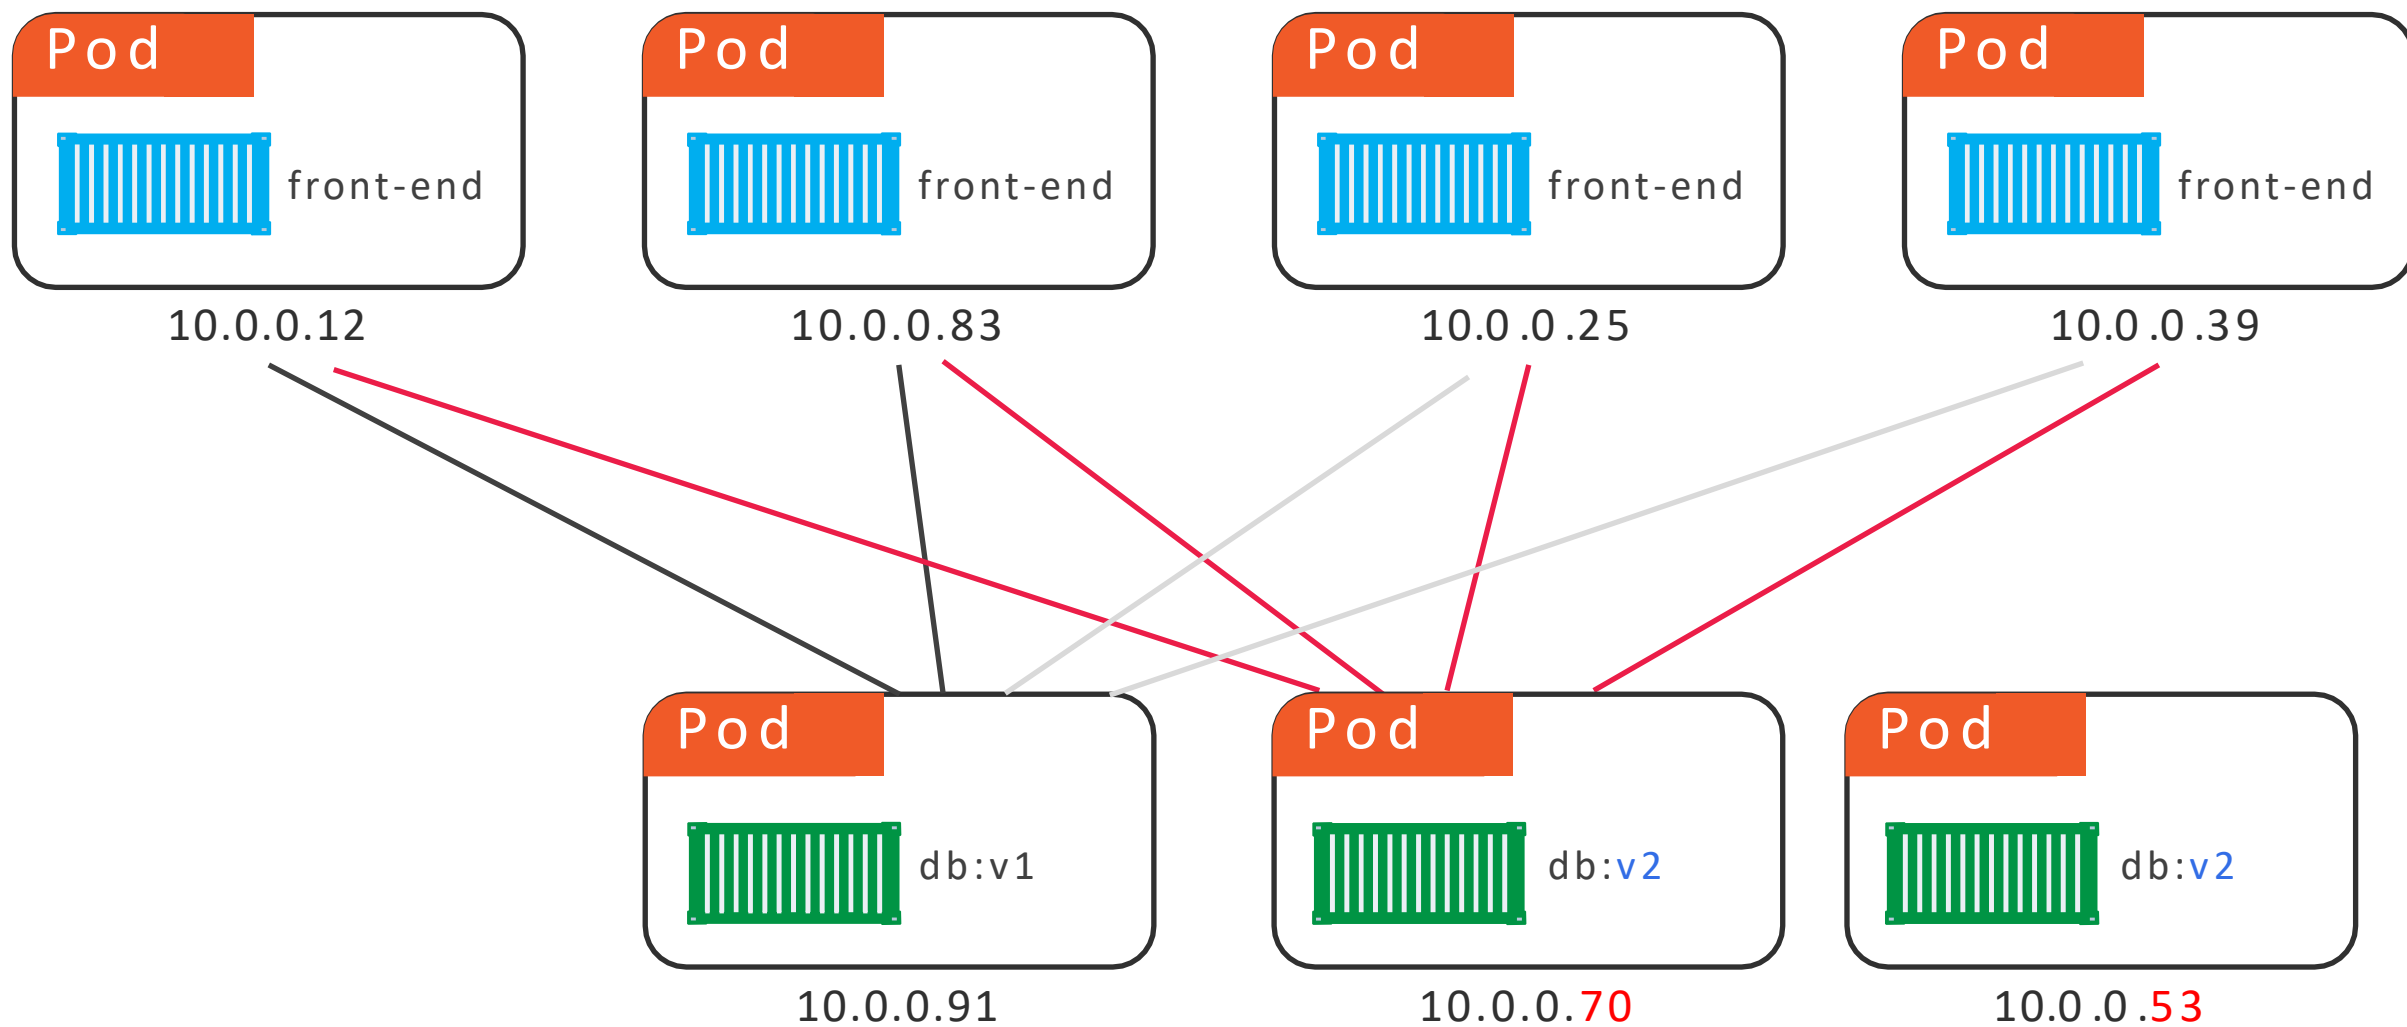
# Services

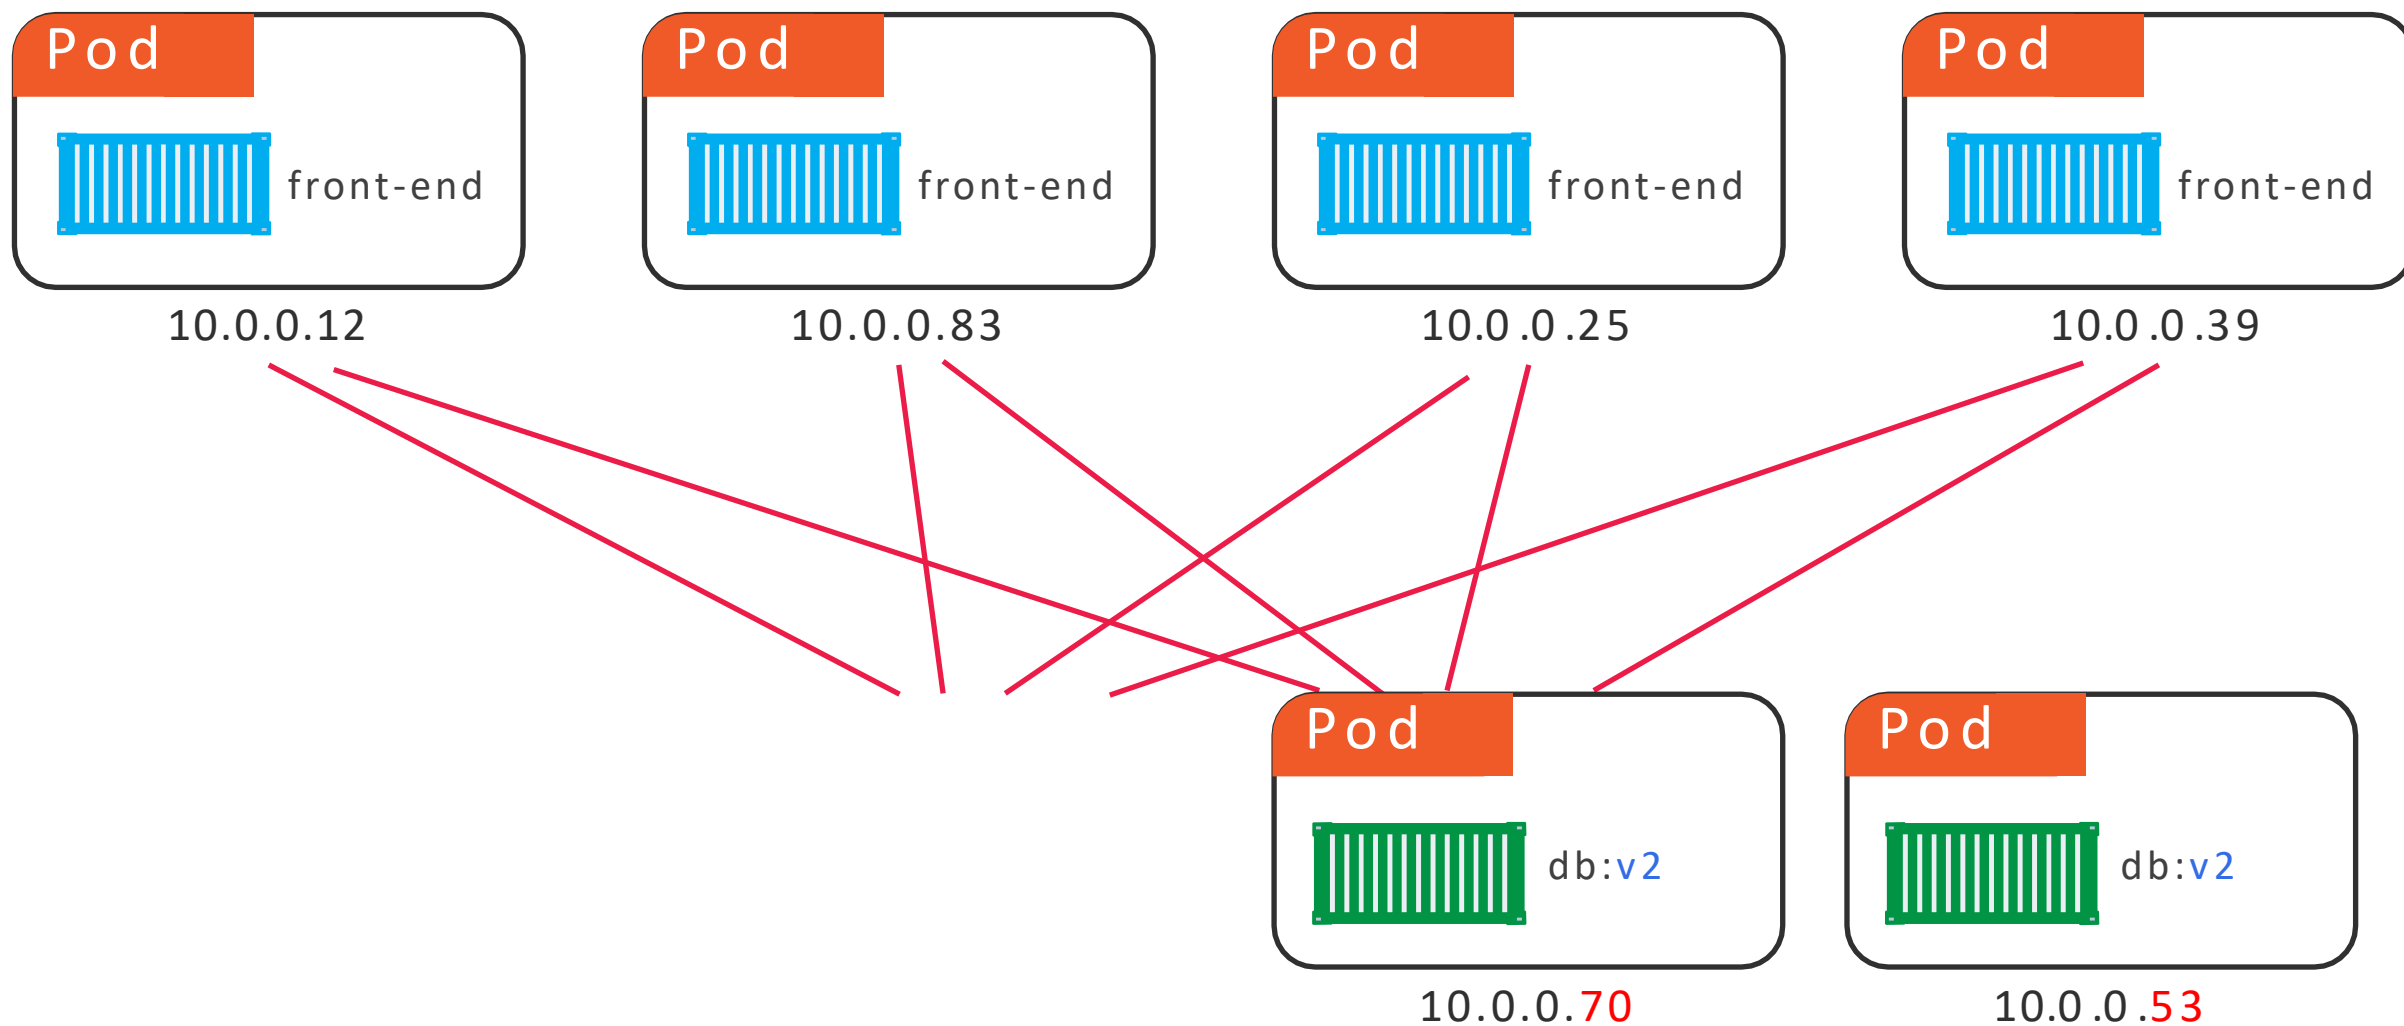


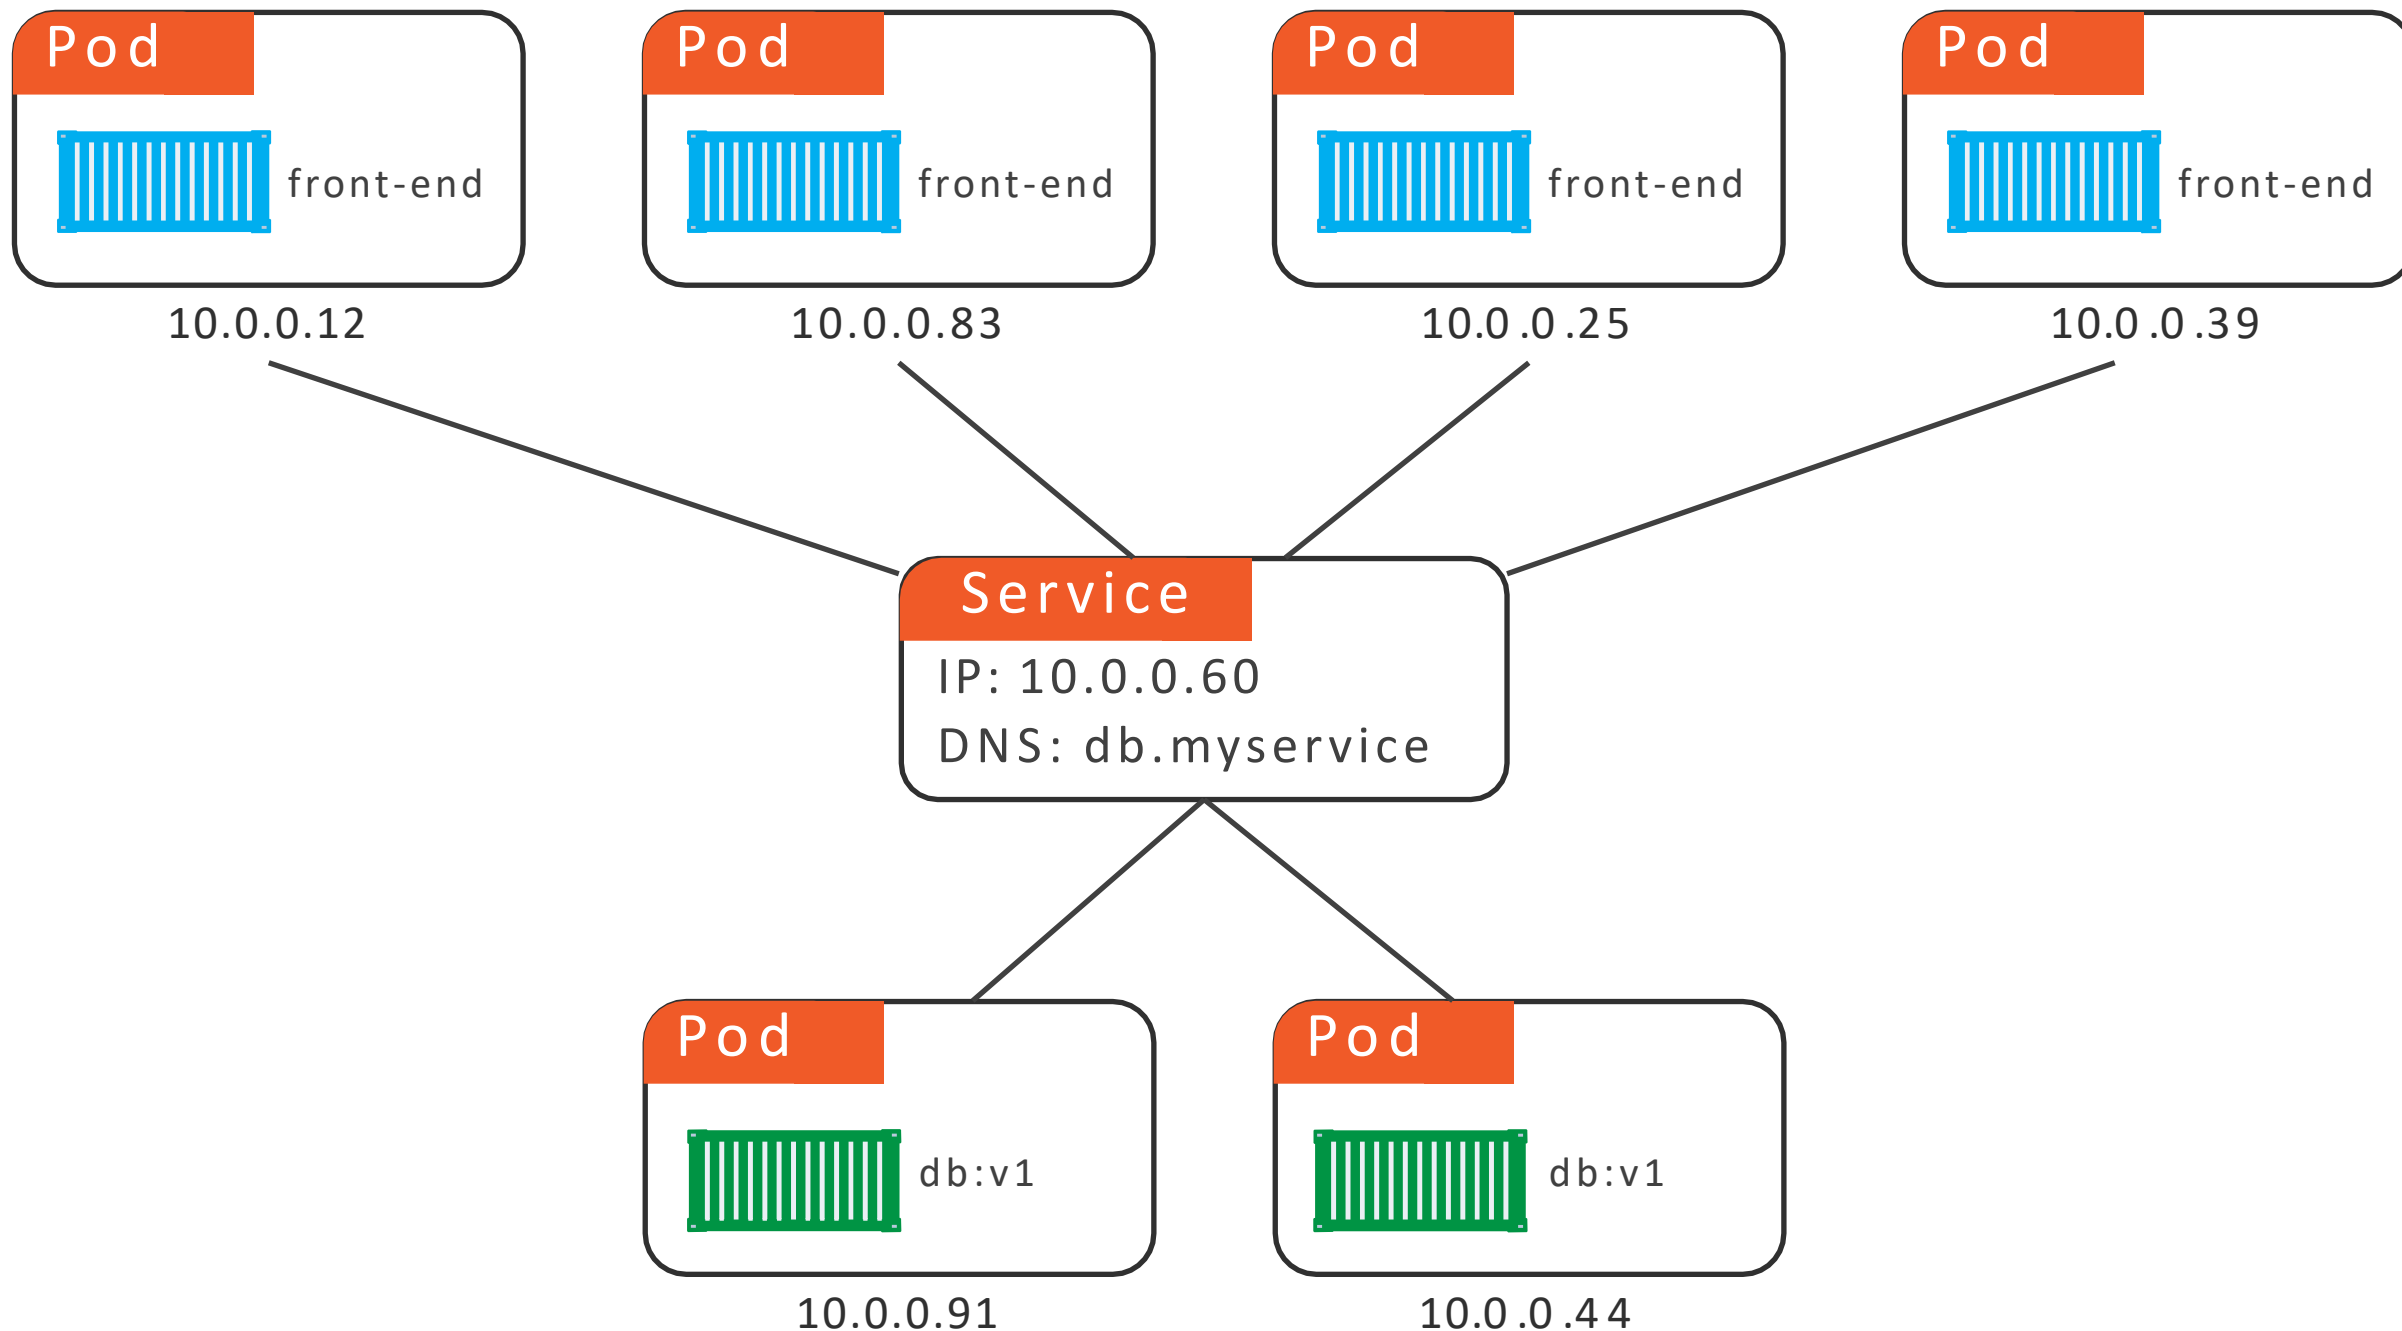
Every new pod gets a new IP = IP churn!

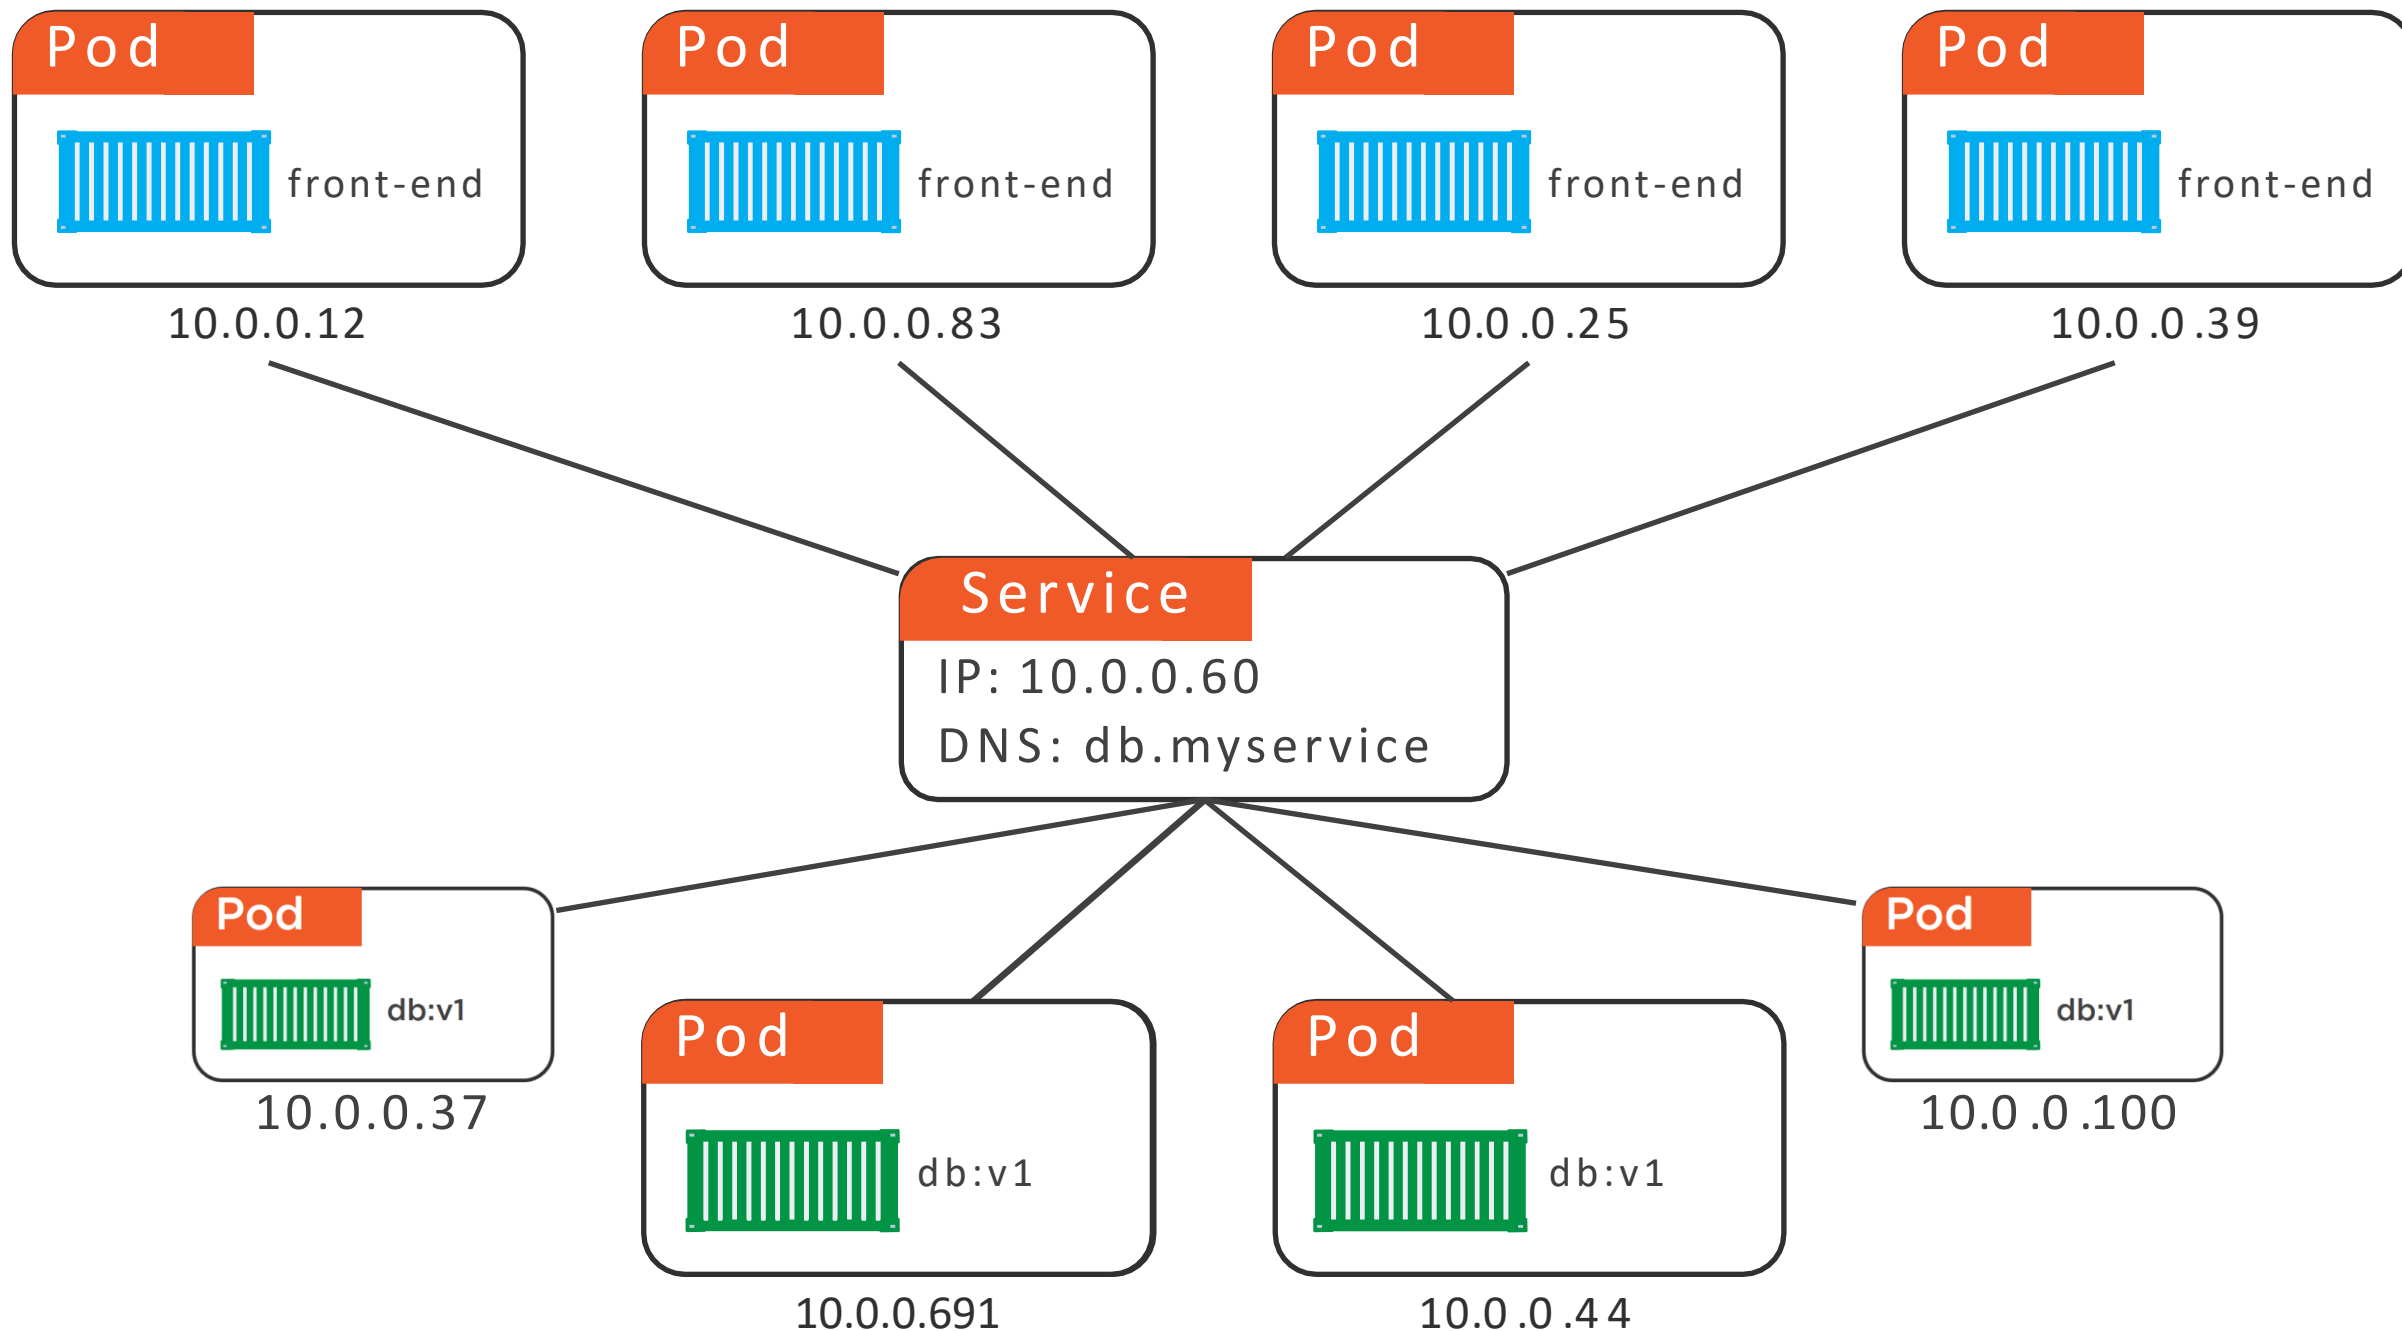


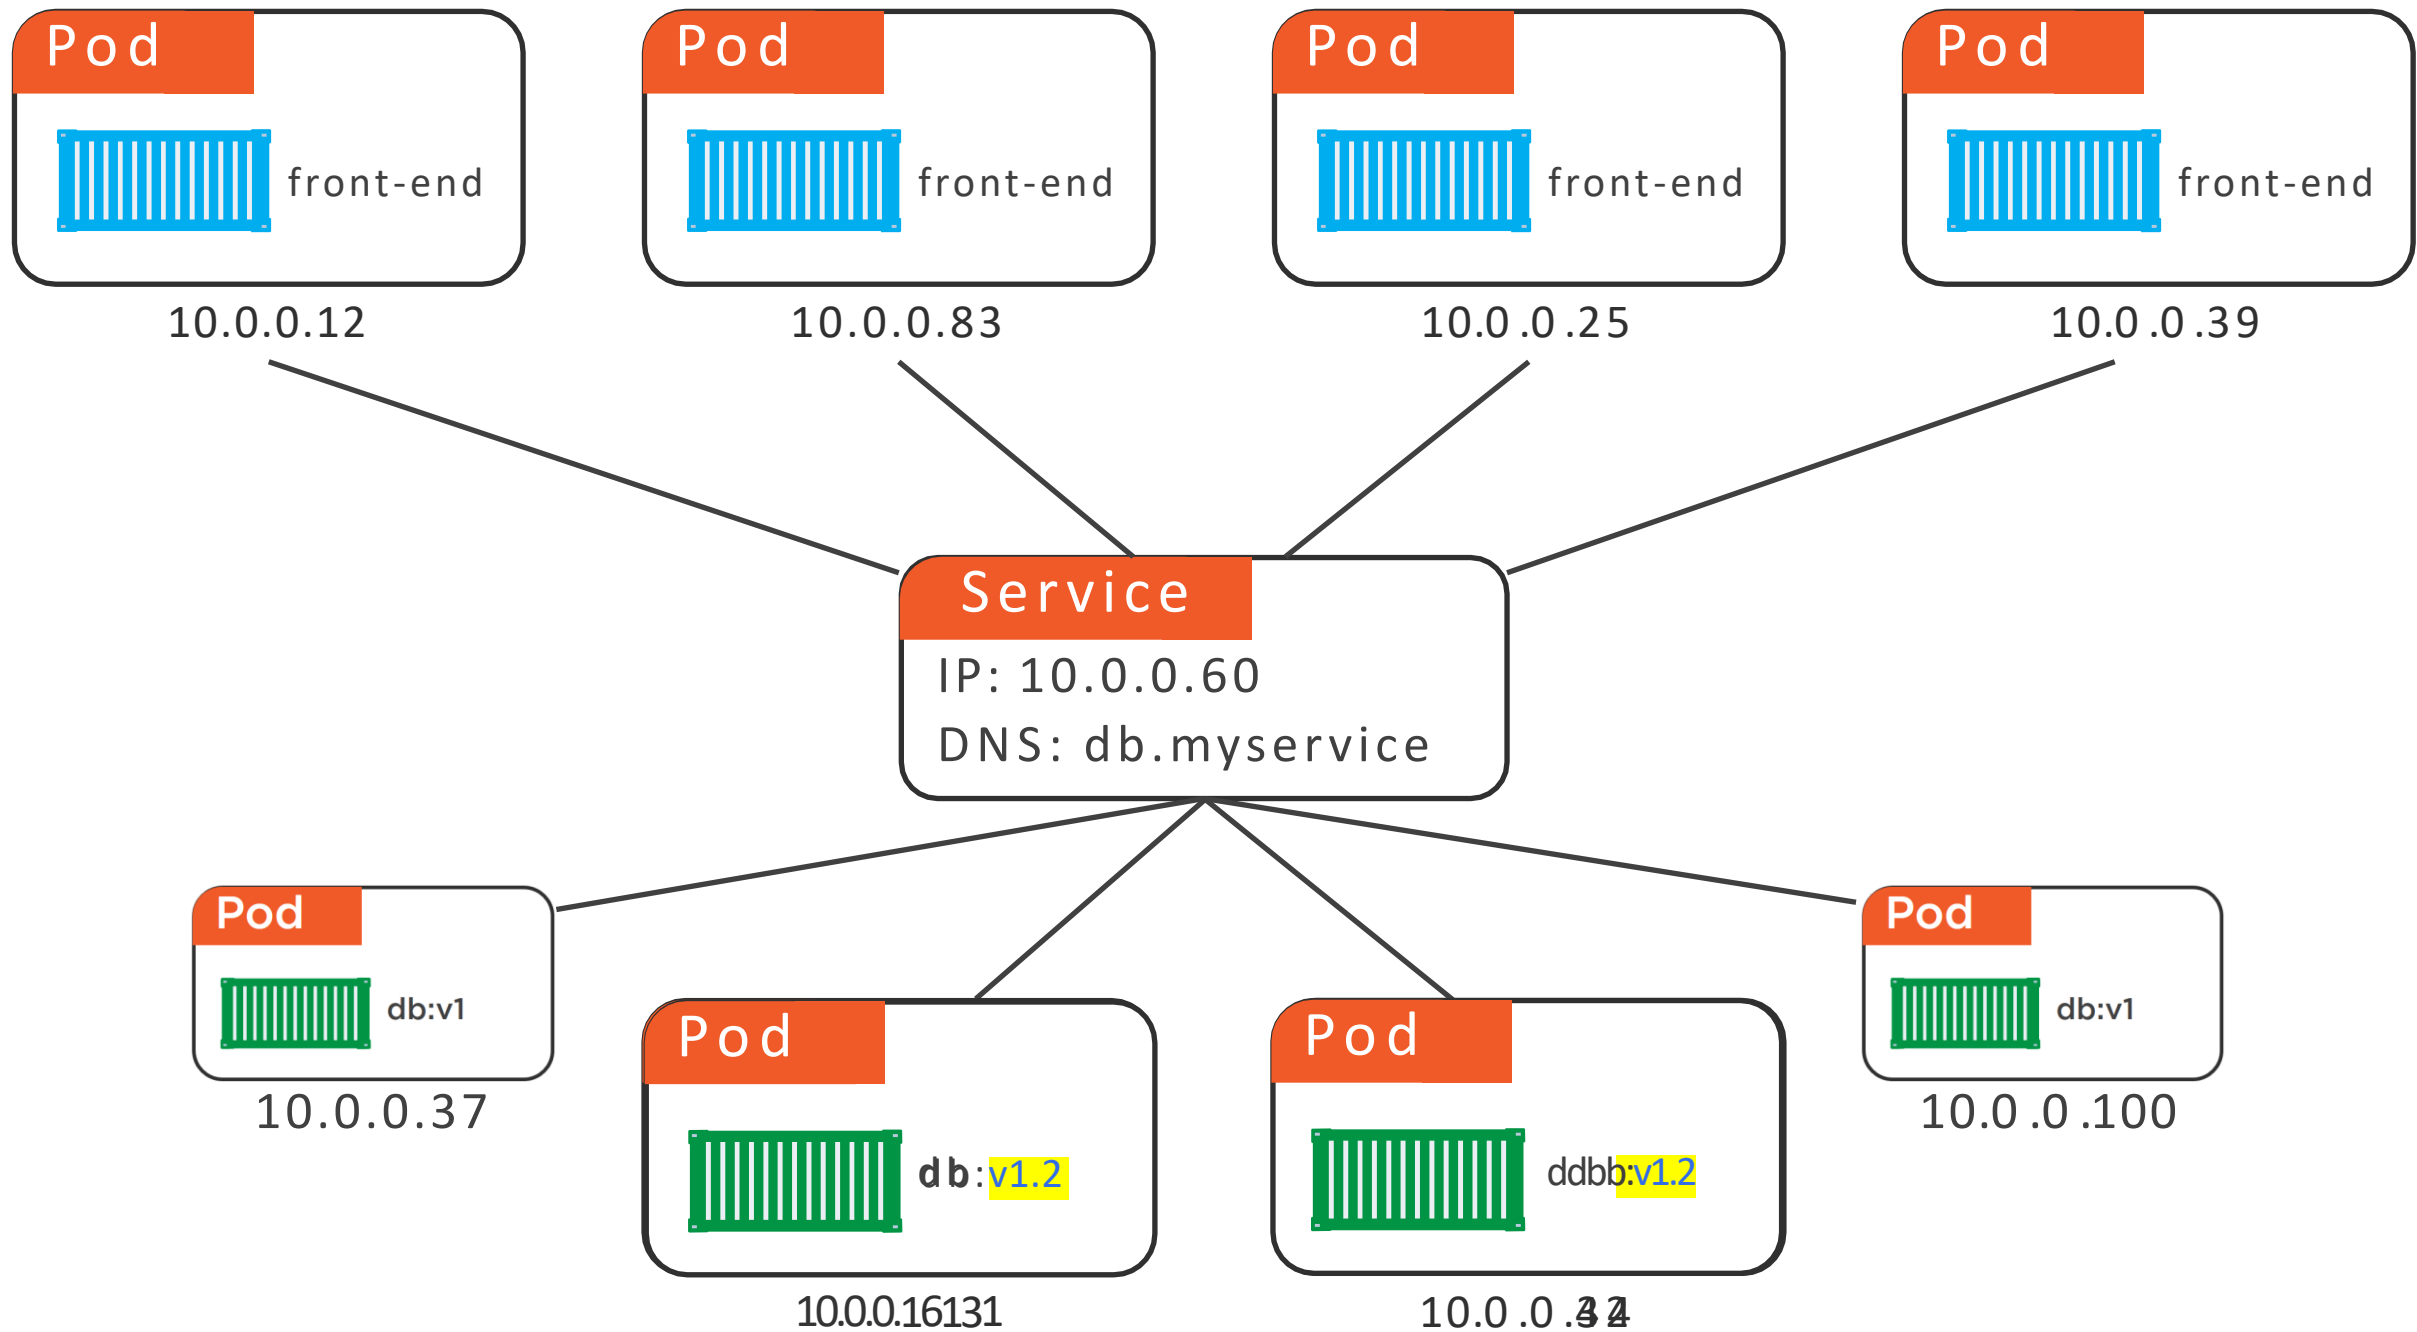












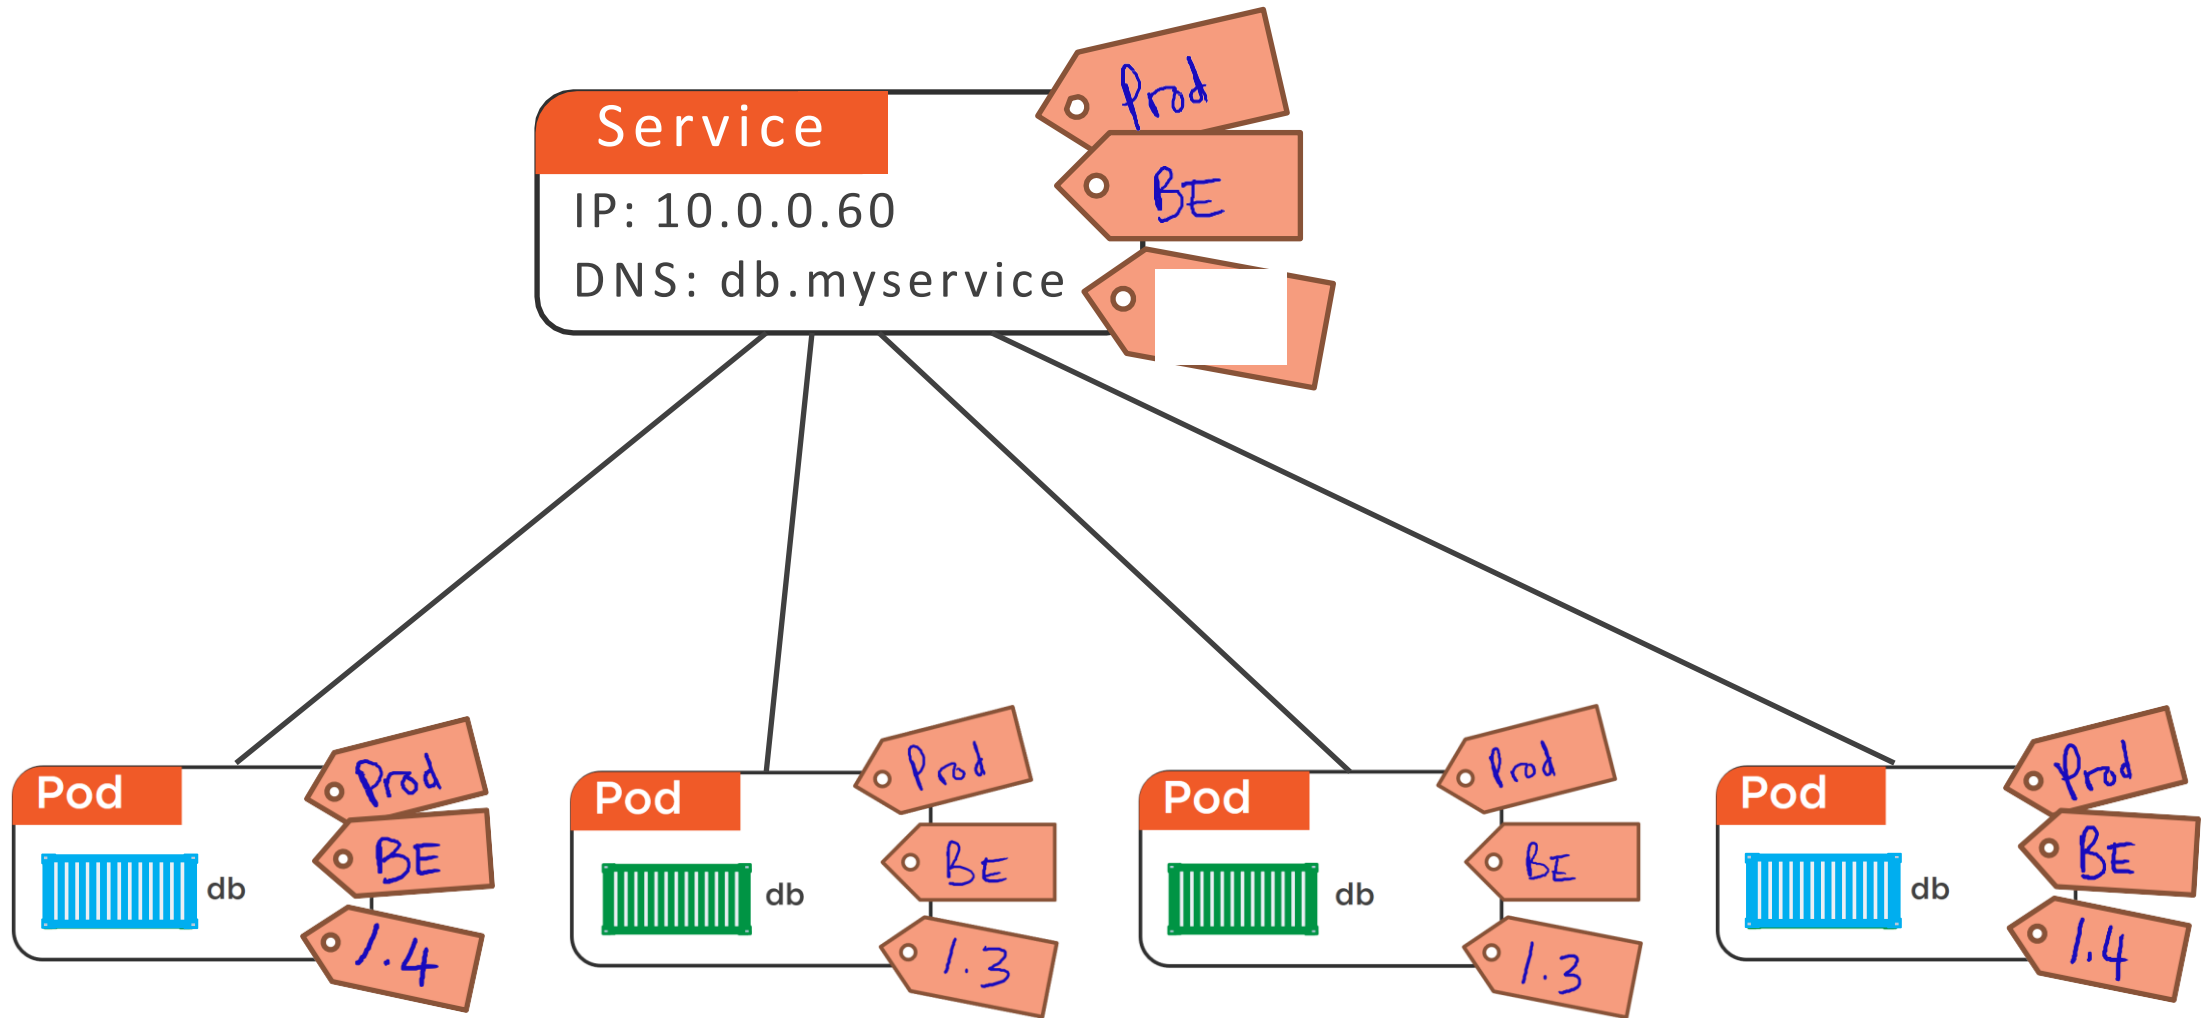




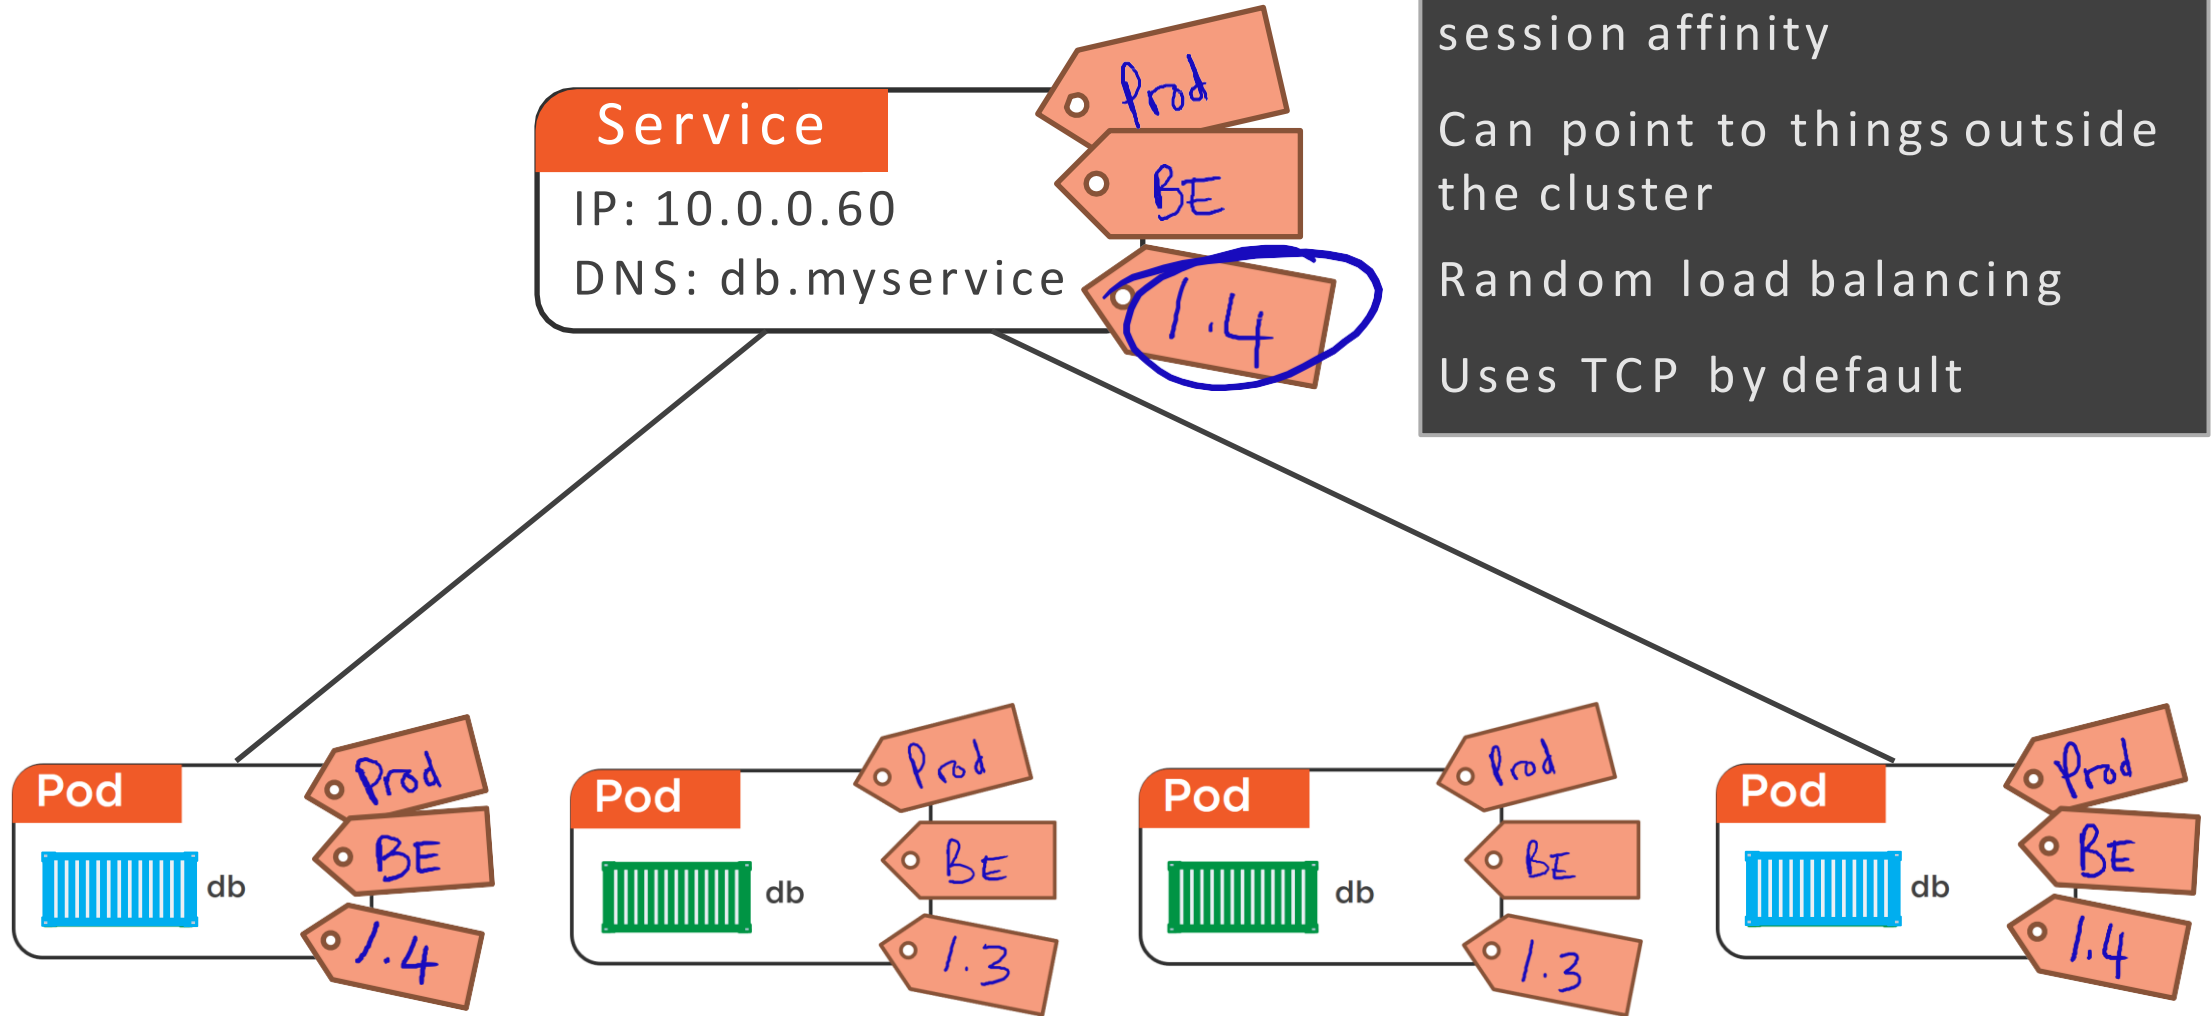










Only send to healthy pods  
Can be configured for session affinity  
Can point to things outside the cluster  
Random load balancing  
Uses TCP by default





# Deployments

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: xyz
spec:
  replicas: 4
```


 Master









apiserver {}


 Node


Pod

 xyz


 Node


Pod

 xyz


 Node

Pod


 xyz


 Node


Pod


 xyz

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: xyz
spec:
  replicas: 4
```


 Master




 KV




apiserver {}


 Node


Pod

 xyz


 Node


Pod

 xyz


 Node

Pod





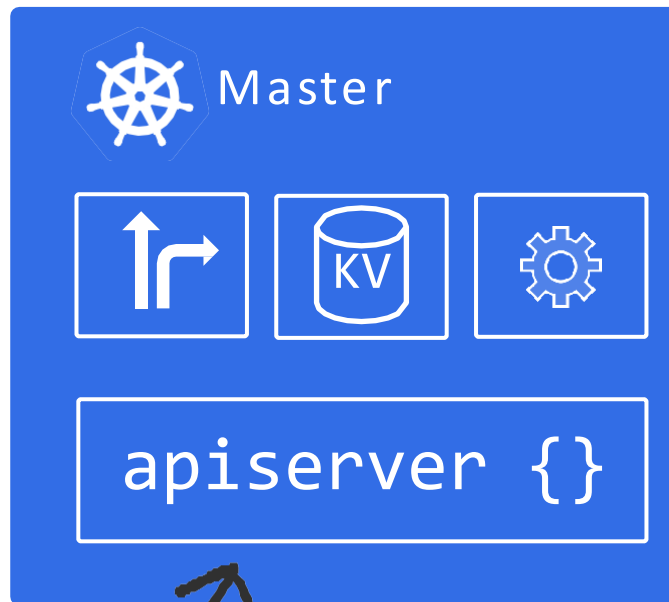
 xyz

 Node

Pod

 xyz

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: xyz
spec:
  replicas: 4
```



REST objects

Self documenting

Deployed via YAML or  
JSON manifests

Spec-once deploy-many

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: xyz
spec:
  replicas: 4
```

Simple rolling updates  
and rollbacks

Add features to  
Replication Controllers  
(Replica Sets) *RC v2*

Versioned

Deployed via the  
apiserver

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: xyz
spec:
  replicas: 4
```






Simple rolling updates  
and rollbacks






Multiple concurrent versions






- Blue-green deployments
- Canary releases





Simple versioned rollbacks













## Cluster Store

- Cluster state and config
- Stateful




## apiserver

- Front-end to control plane






Kubelet  
Main Kubernetes agent



Container engine  
Docker or rkt



kube-proxy  
Kubernetes networking



Objects  
in the  
K8s API



Pods :Atomic unit of scheduling...

Replication

Controllers :Scale pods, desired state etc...

Deployments :RC +rolling updates, rollbacks...

Services :Stable networking...

Coming up next...

# Installing Kubernetes