

# Working with Pods

---





Theory


Hands-on

# Theory

Hypervisor  
Virtualization

VM

Atomic unit of  
scheduling



Container


Atomic unit of  
scheduling




Pod

Atomic unit of  
scheduling


# Theory





```
apiVersion: v1
kind: Pod
metadata:
  name: hello-pod
spec:
  containers:
    - name: hello-ctr
      image:...
      ports:
        - containerPort: 8080
```





# Theory



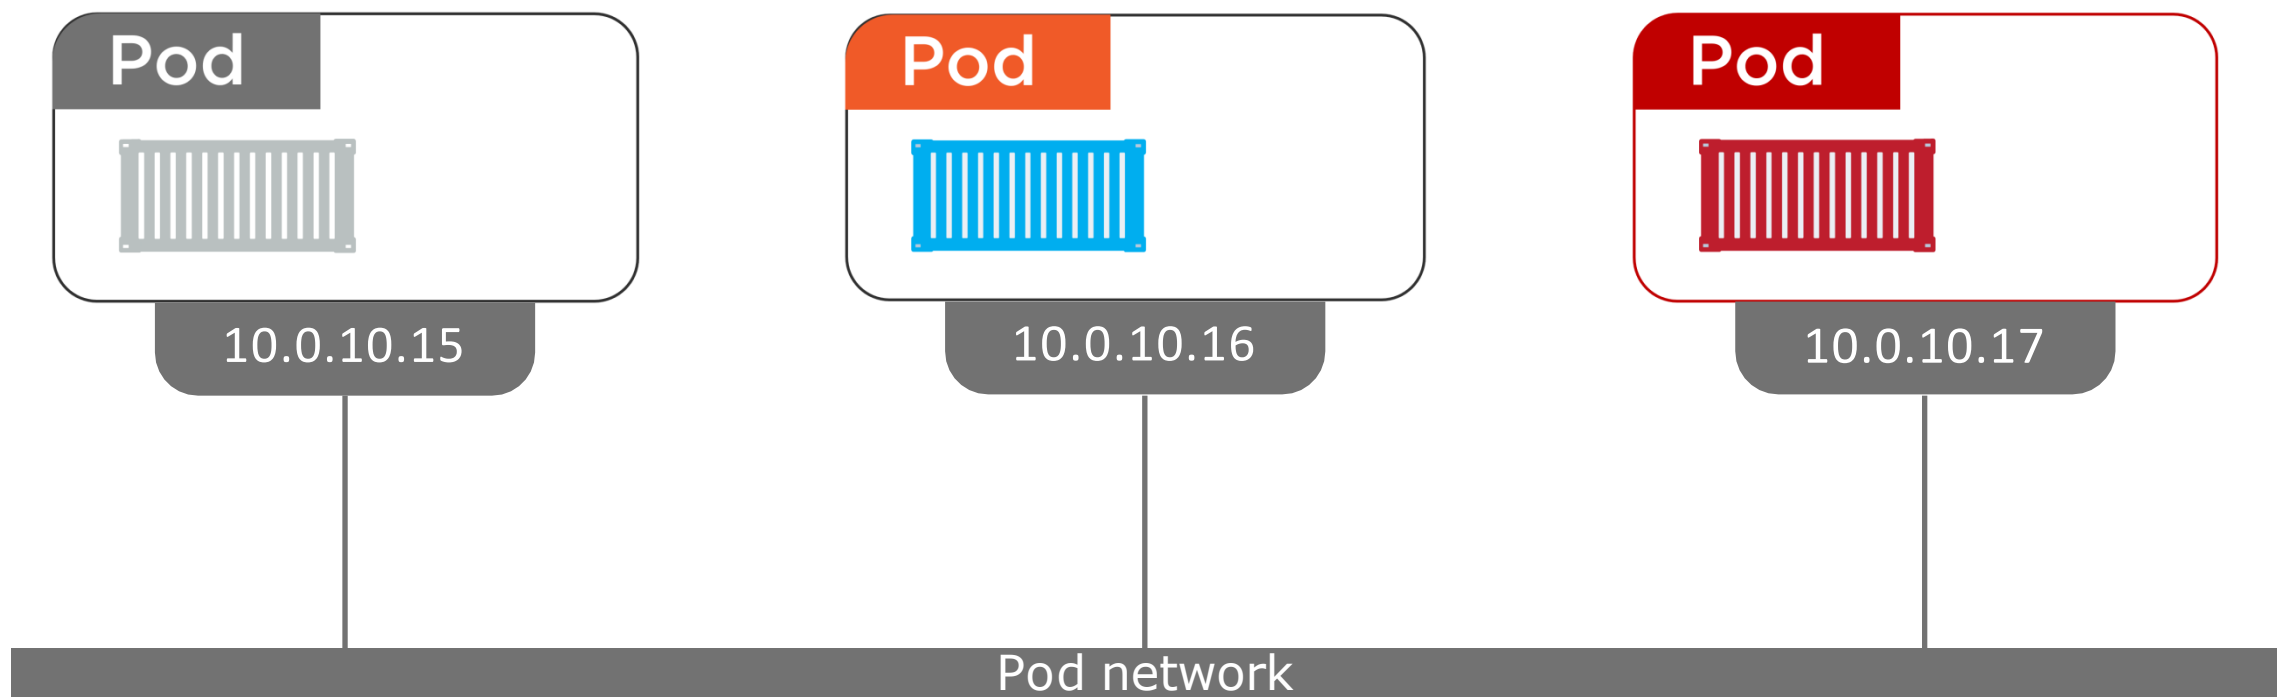
```
apiVersion: v1
kind: Pod
metadata:
  name: hello-pod
spec:
  containers:
    - name: hello-ctr
      image:...
      ports:
        - containerPort: 8080
```




# Theory



# Inter-pod Communication




# Inter-pod Communication








# Intra-pod Communication





# Intra-pod Communication





# Pods: Atomic Deployments





# Pods: Atomic Deployments



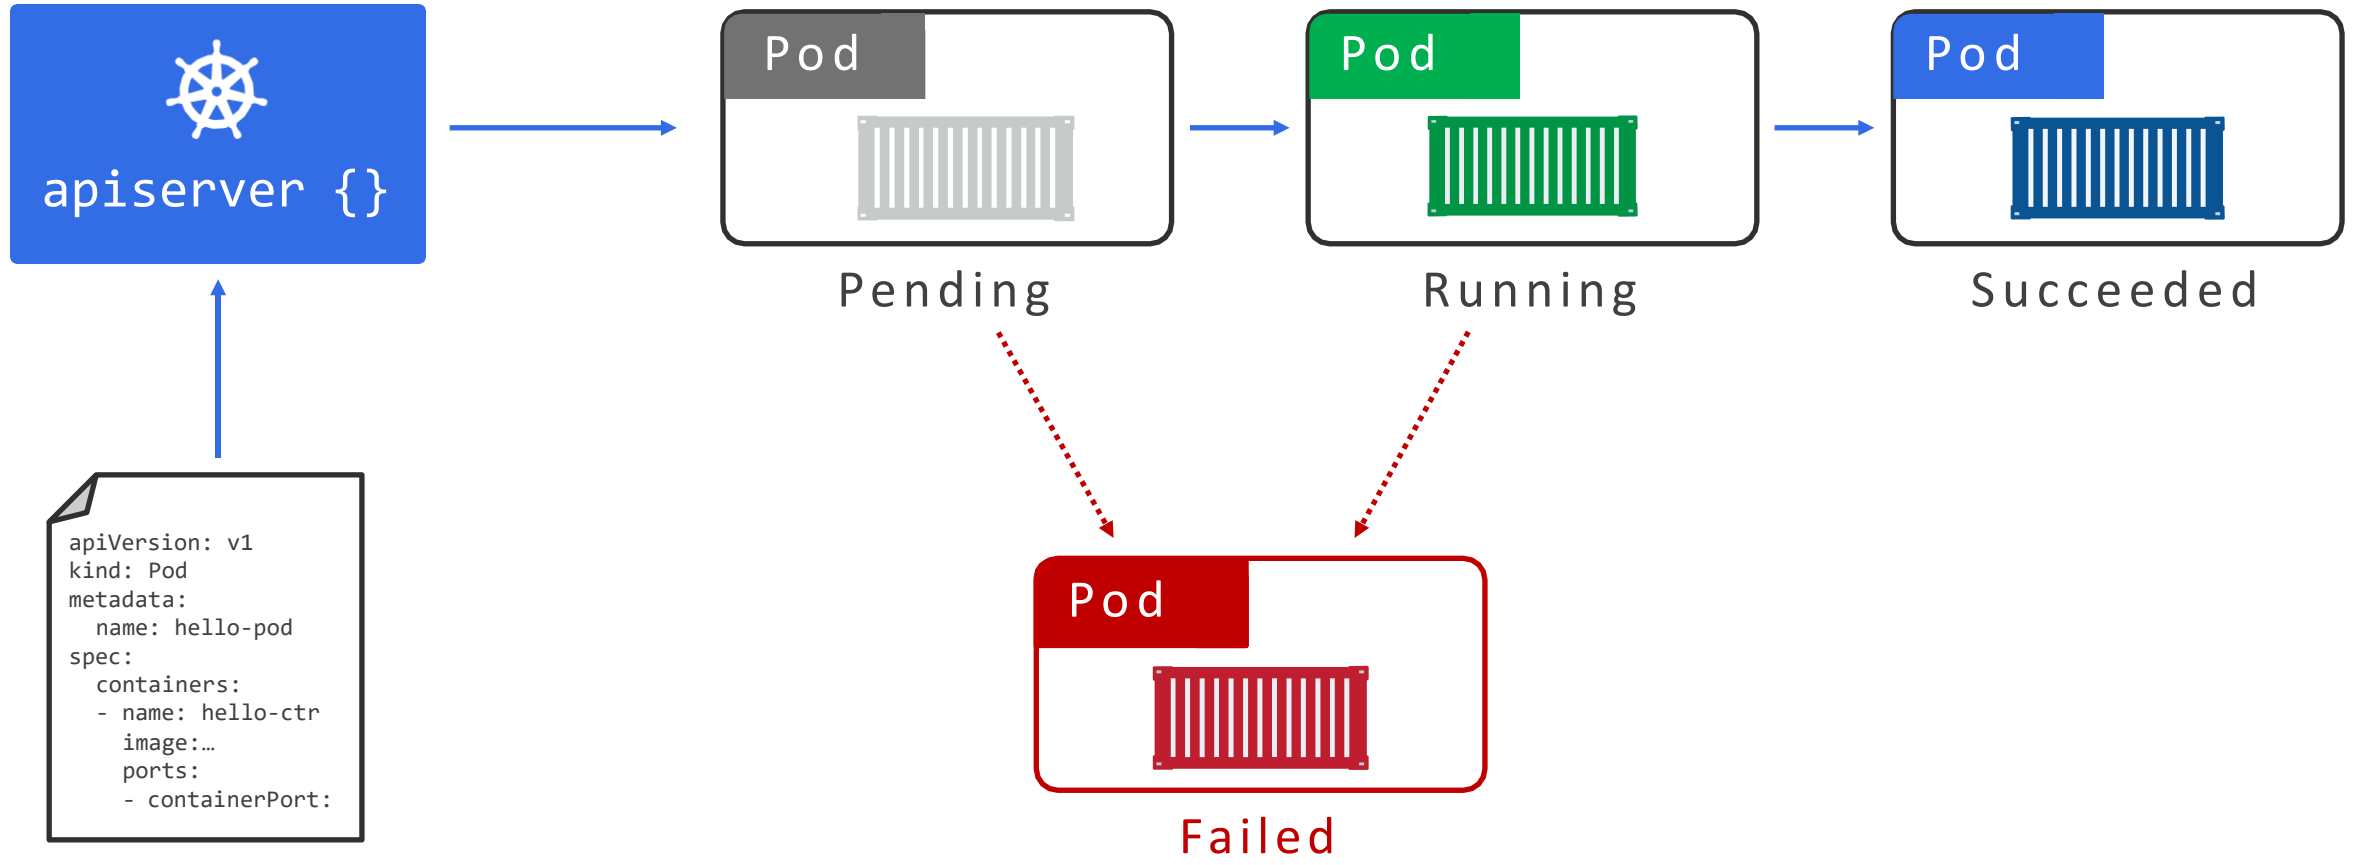
# Pods: Atomic Deployments



# Pods:



# Pods:



# Pods:

Smallest/atomic unit  
of scheduling

Have one or more  
containers

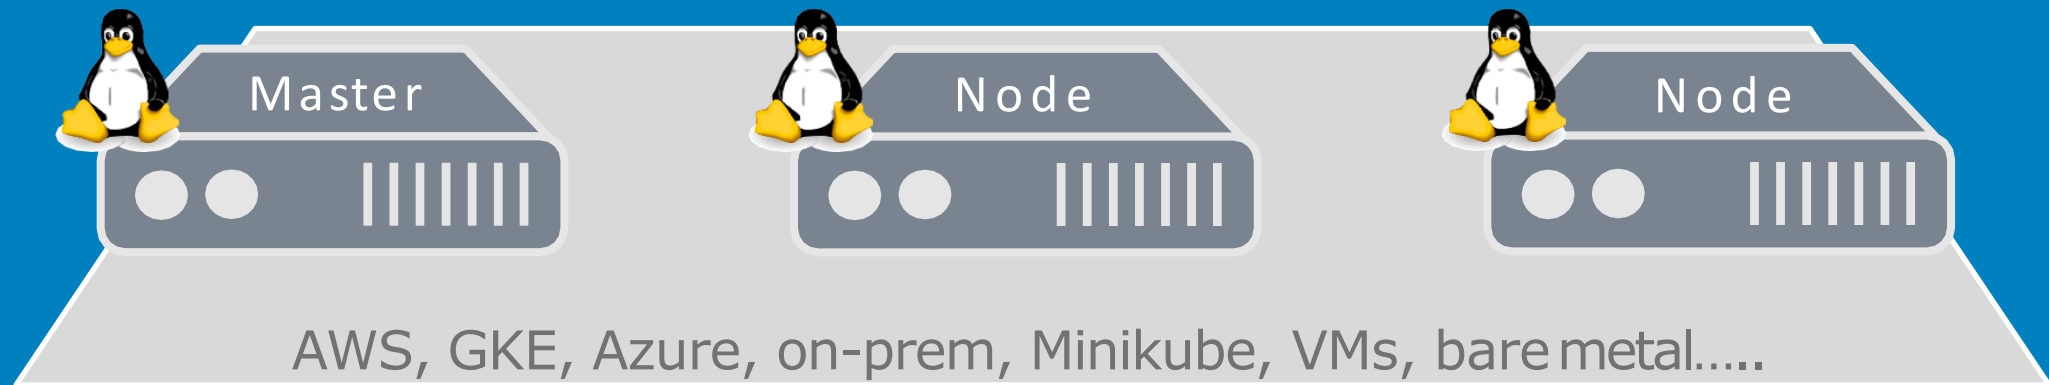


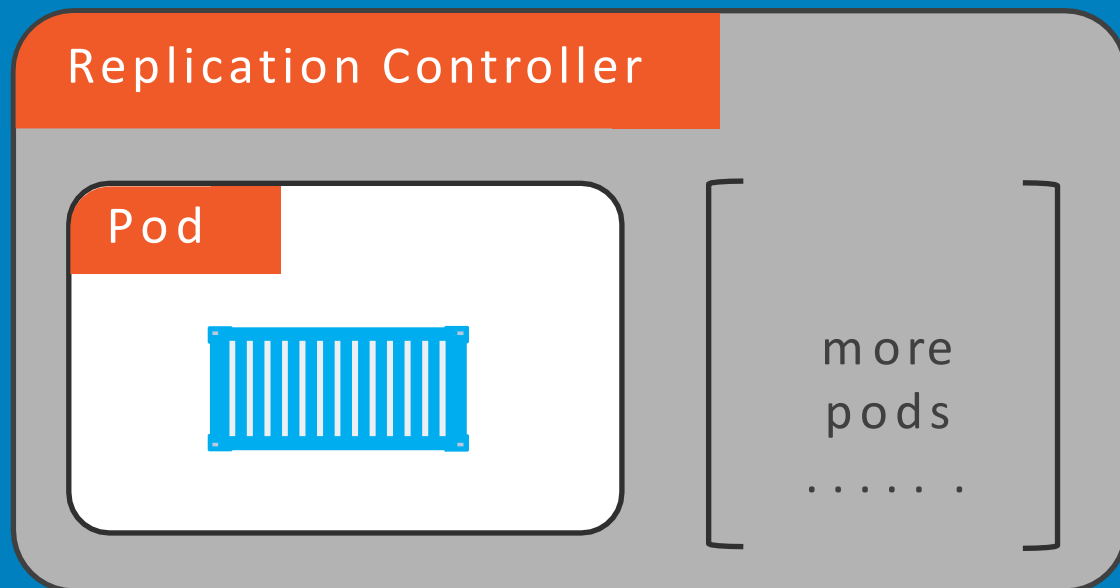

Scheduled on nodes  
(minions)

Declarative via  
manifest files



```
$ kubectl
```





# Summary

Smallest unit of  
scheduling  
(Atomic)

Pod manifests  
(YAML or JSON)




Deployed via other  
objects  
(Replication Controllers etc...)

Replication  
Controllers  
(Implement desired state)

Deployment

Replication Controller

Pod



more  
pods  
.....